

# Overview & I/O Structure

Prof. Li-Pin Chang  
CS@NYCU

# THE THREE PILLARS (of an operating system)

# Pillar 1: Process Management

- A process is a program in execution; Program is a **passive** entity, process is an **active** entity
- Process needs resources to accomplish its task
  - CPU time for execution
  - Memory space for program store
  - I/O bandwidth for communication
  - Storage space for data
- A system has many processes running concurrently on one or more CPUs

# Process Management Activities

- The operating system is responsible to the following activities regarding process management:
  - Creating and terminating processes
  - Providing mechanisms for process communication
  - Providing mechanisms for process synchronization

# Pillar 2: Memory Management

- All **data** and **instructions** must be in memory for read/write and execution, respectively
- Memory management activities
  - Allocating and deallocating memory space as needed
  - Deciding which processes (or parts thereof) and data to move into and out of memory (**paging**)
  - Allowing programs allocate much more memory than physical memory (**virtual memory**)
  - Keeping track of which parts of memory are currently being used and by whom (**protection**)

# Pillar 3: File/Storage Management

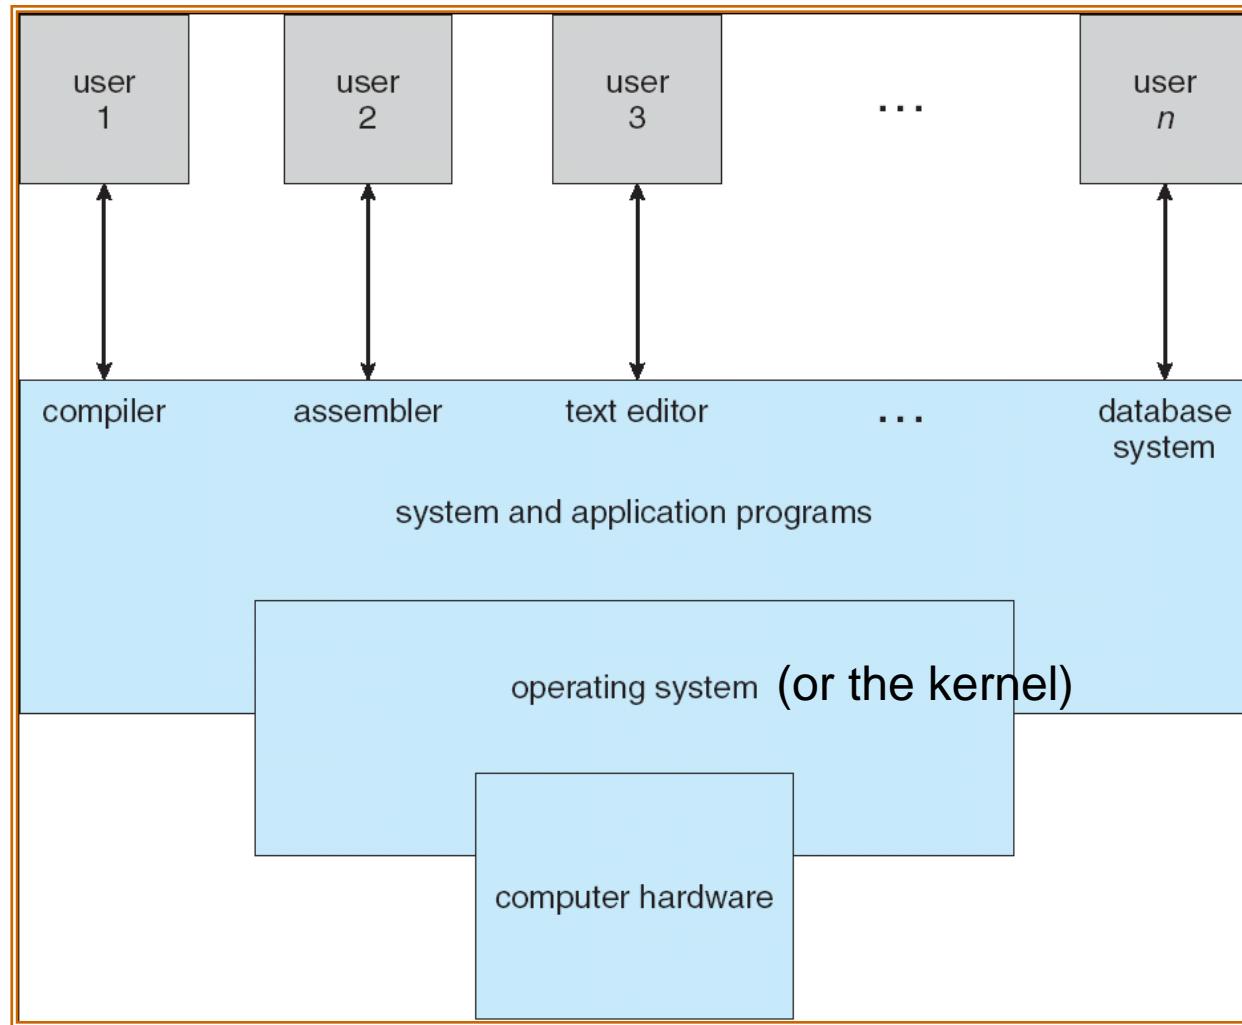
- OS provides uniform, logical view of information storage, i.e., blocks and file
- File-System management
  - File and directory organization
  - Access control on to determine who can access what
  - Storage space management
- File system basic activities
  - Creating and deleting files and directories
  - Reading and writing files and directories
  - Mapping files from secondary storage to main memory
  - Security and protection
  - Crash recovery and data consistency

# Storage Management

- Entire speed of computer operation hinges on storage subsystem and its algorithms
  - Storage ~1 ms; memory ~100ns;
- Optimizing latency and throughput of data access
  - I/O scheduling
  - Data caching and buffering

# OPERATING SYSTEM DEFINITION

# Computer System Overview



# What is an Operating System?

- A program that acts as an intermediary between a user of a computer and the computer hardware
- Operating system goals:
  - Execute user programs and make solving user problems easier
  - Make the computer system convenient to use
  - Use the computer hardware in an efficient manner

# Operating System Definition

- OS is a resource allocator
  - Manages all resources
  - Decides between conflicting requests for **efficient** and **fair** resource use
- OS is a control program
  - Controls (or monitor) execution of programs to **prevent errors and improper use** of the computer

# Operating System Definition (Cont'd)

No black-and-white definition

- “*Everything a vendor ships when you order an operating system*” is a rough approximation
  - Retrospect: Internet Explorer in Windows 95
- “The one program running at all times on the computer” is the **kernel**
  - Everything else is either a user application or system program
- Example: the UNIX OS
  - POSIX (IEEE 1003.1-2017)
  - A kernel, plus
  - A collection of system programs  $\sim=$  coreutil + binutil

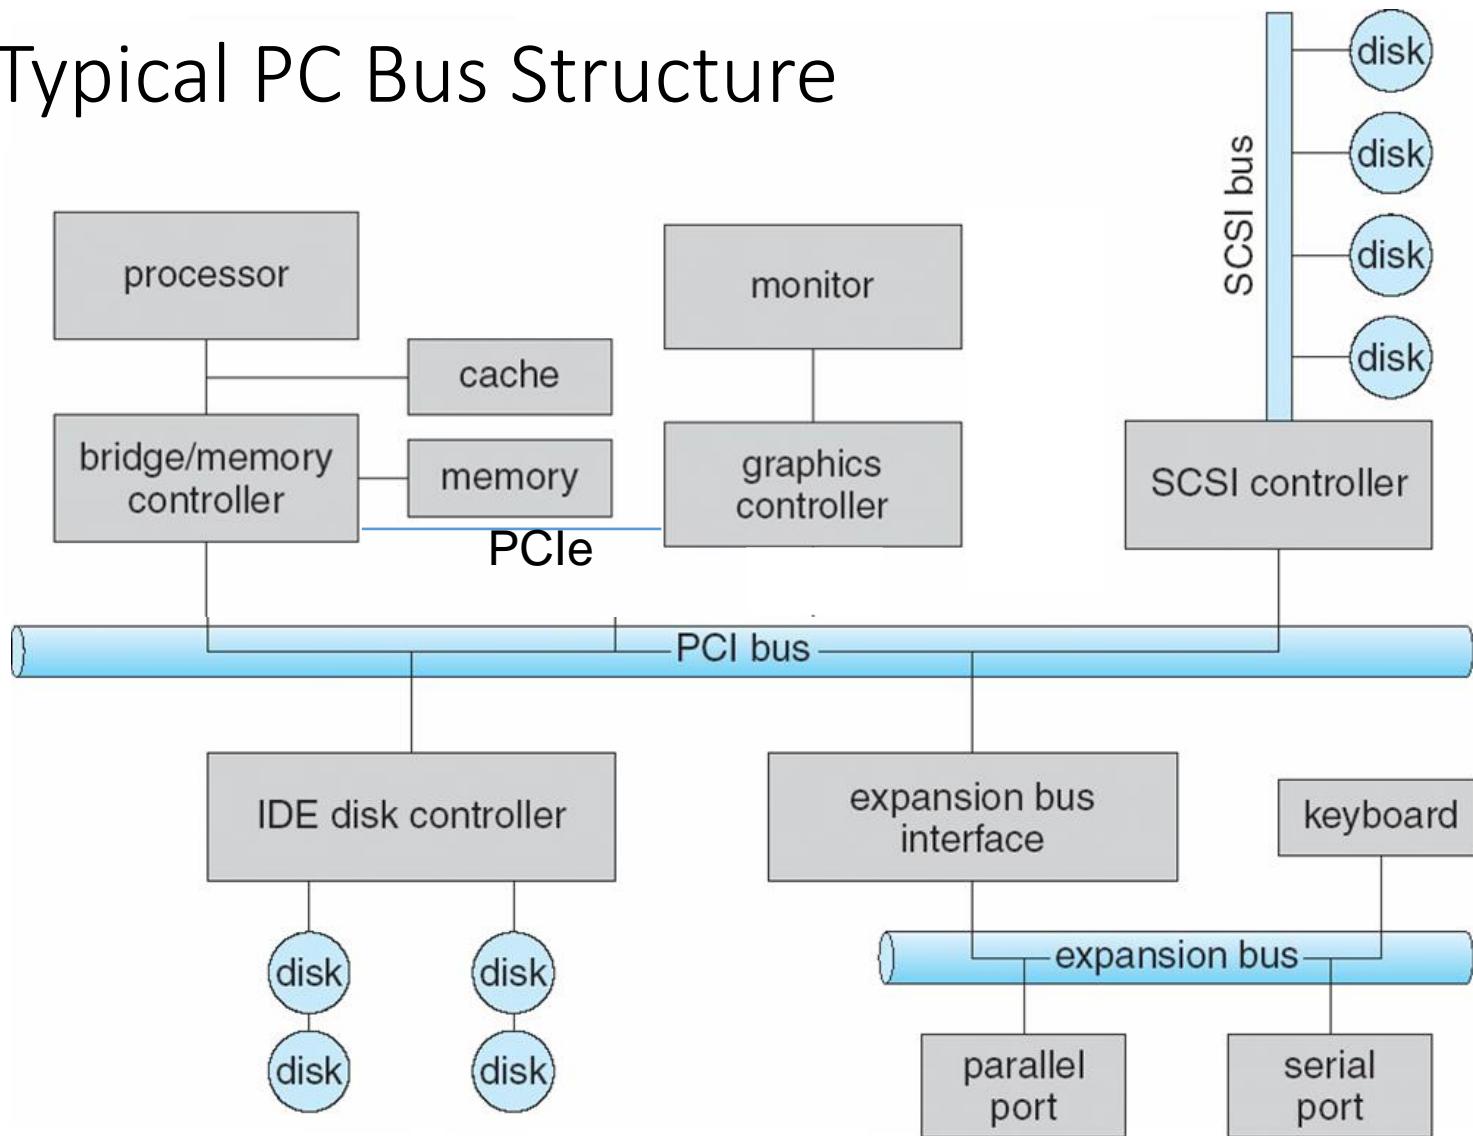
- Which one(s) of the following are part of operating systems?
  - a) The CPU scheduler
  - b) Compilers
  - c) Word processors
  - d) Binary utilities, such as ls, cp, etc

# I/O STRUCTURE

# Why Bothering with I/Os?

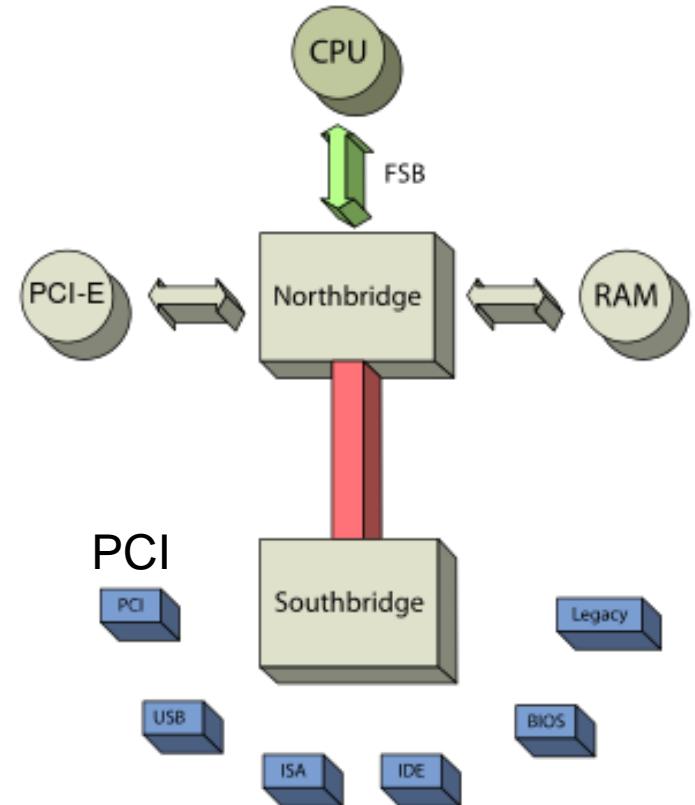
- The kernel is built surrounding interrupts
- The sequential execution of programs is just an illusion that the OS creates; in fact, the CPU jumps back and forth among user programs (multitasking)
- User programs run on the CPU in parallel with operations on I/O devices; OS strives to optimize CPU utilization and I/O utilization at the same time

# A Typical PC Bus Structure



# Typical PC Organization

- North bridge (memory hub)
  - CPU, Memory, fast PCIe devices, like graphics, NVMe SSDs
  - Has been integrated into the CPU
- South bridge (IO hub)
  - PCI, USB, SCSI, SATA
  - ...



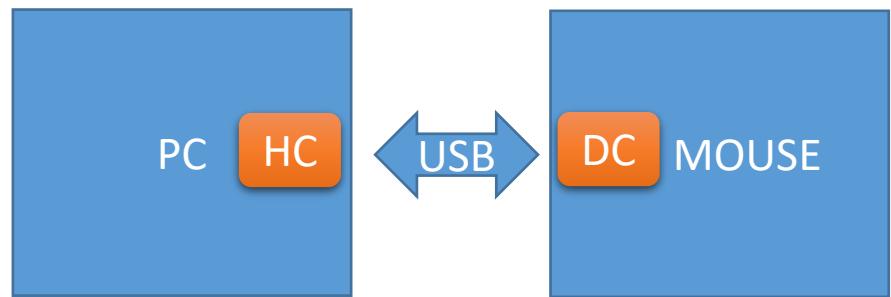
Remark

PCI (parallel, slow, on south)

PCIe (serial, fast, on north)

# I/O Hardware

- Incredible variety of I/O devices
- Common concepts
  - Host controller
  - Bus or interface
  - Device controller
- I/O instructions control devices
  - Direct I/O
  - Memory-mapped I/O



# Partial Device I/O Ports for Legacy PCs (Direct I/O)

I/O address range (hexadecimal)	device
000–00F	DMA controller
020–021	interrupt controller
040–043	timer
200–20F	game controller
2F8–2FF	serial port (secondary)
320–32F	hard-disk controller
378–37F	parallel port
3D0–3DF	graphics controller
3F0–3F7	diskette-drive controller
3F8–3FF	serial port (primary)

# Example: Set the x86 Interval Timer (Direct I/O)

```
mov al, 0x36
out 0x43, al      ;tell the PIT which channel we're setting

mov ax, 11931
out 0x40, al      ;send Low byte          11931820 / x = f.
mov al, ah          Set x=11931 → f=1000Hz
out 0x40, al      ;send high byte
```

Channel 0	0x40	System Timer
Channel 1	0x41	DRAM refresh (obsolete)
Channel 2	0x42	Buzzer
Command	0x43	Command I/O register

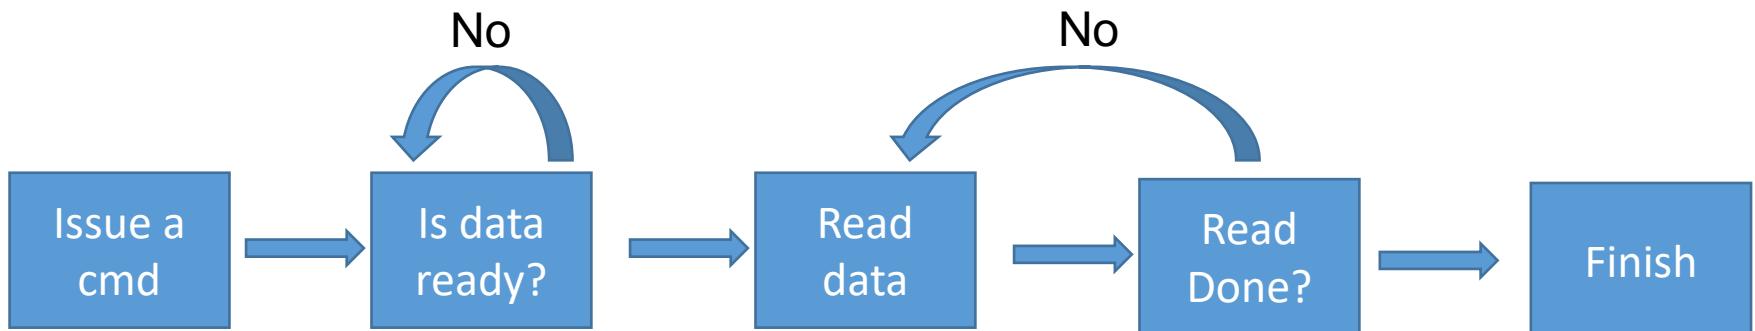
# I/O Operations

I/O devices and CPUs execute concurrently

- **Polling** I/O: The CPU waits on an I/O device until completion
- **Interrupt** I/O: The CPU will be notified of I/O completion

# Polling

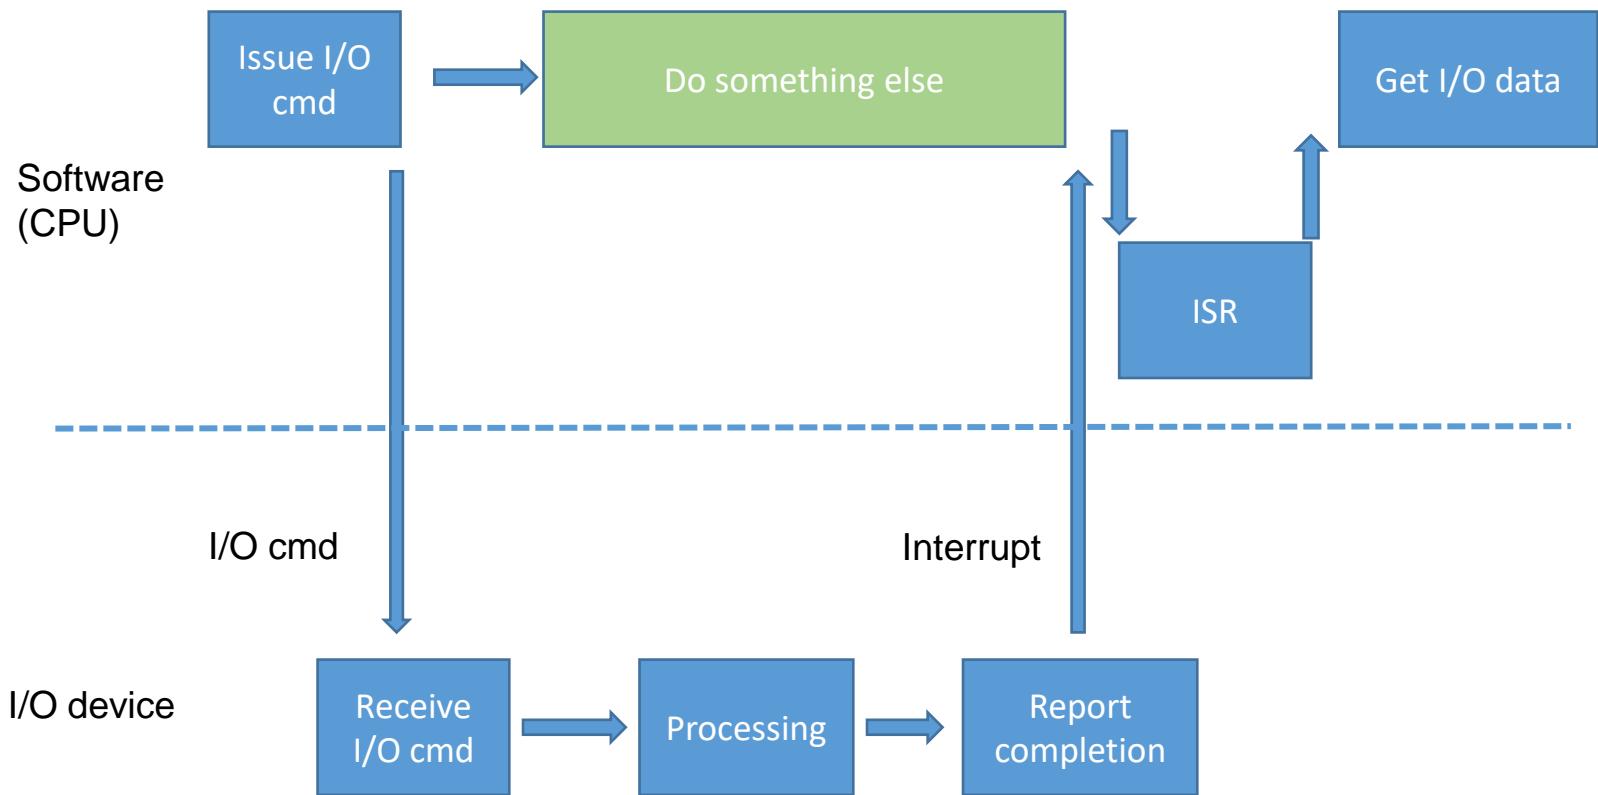
- CPU determines state of device
  - Busy, data ready, or error
- Busy-wait cycle to wait for I/O from device
- CPU performs data movement
- For example, GPIO



# Interrupts

- CPU Interrupt-request line triggered by I/O device
- Interrupt vector to dispatch interrupt to correct handler
- Interrupt handler (ISR) receives interrupts
- Interrupt mechanism also used for exceptions

# A (Simplified) Interrupt-Driven I/O



# Common Functions of Interrupts

- Interrupt transfers control to the interrupt service routine (ISR) generally, through the **interrupt vector**, which contains the addresses of all the service routines
  - A **trap** is a software-generated interrupt caused either by an error or a user request
  - An **IRQ** is generated by hardware
- An operating system is interrupt-driven
- **Hardware:** interrupt request (IRQ)
  - I/O completion, timer, etc
- **Software:** Trap
  - divide by zero, access violation, system call

# Legacy PC interrupts

Interrupt	Function
00h	Divide By Zero
01h	Single Step
02h	Nonmaskable Interrupt (NMI)
03h	Breakpoint Instruction
04h	Overflow Instruction
05h	Print Screen
05h	Bounds Exception (80286, 80386)
06h	Invalid Op Code (80286, 80386)
07h	Math Coprocessor Not Present
08h	Double Exception Error (80286, 80386) (AT Only)
08h	System Timer - IRQ 0
09h	Keyboard - IRQ 1
09h	Math Coprocessor Segment Overrun (80286, 80386) (AT Only)
0Ah	IRQ 2 - Cascade from Second programmable Interrupt Controller
0Ah	Invalid Task Segment State (80286, 80286) (AT Only)
0Ah	IRQ 2 - General Adapter Use (PC Only)
0Bh	IRQ 3 - Serial Communications (COM 2)
0Bh	Segment Not Present (80286, 80386)
0Ch	IRQ 4 - Serial Communications (COM 1)
0Ch	Stack Segment Overflow (80286, 80386)
0Dh	Parallel Printer (LPT 2) (AT Only)
0Dh	IRQ 5 - Fixed Disk (XT Only)
0Dh	General Protection Fault (80286, 80386)
0Eh	IRQ 6- Diskette Drive Controller
0Eh	Page Fault (80386 Only)
0Fh	IRQ 7 - Parallel printer (LPT 1)
•	
•	
•	

# Interrupt Handling

- The operating system preserves the state of the CPU by storing **registers** and the **program counter**
- Determines which type of interrupt has occurred:
  - Reading I/O registers
  - vectored interrupt system
- Separate segments of code determine what action should be taken for each type of interrupt

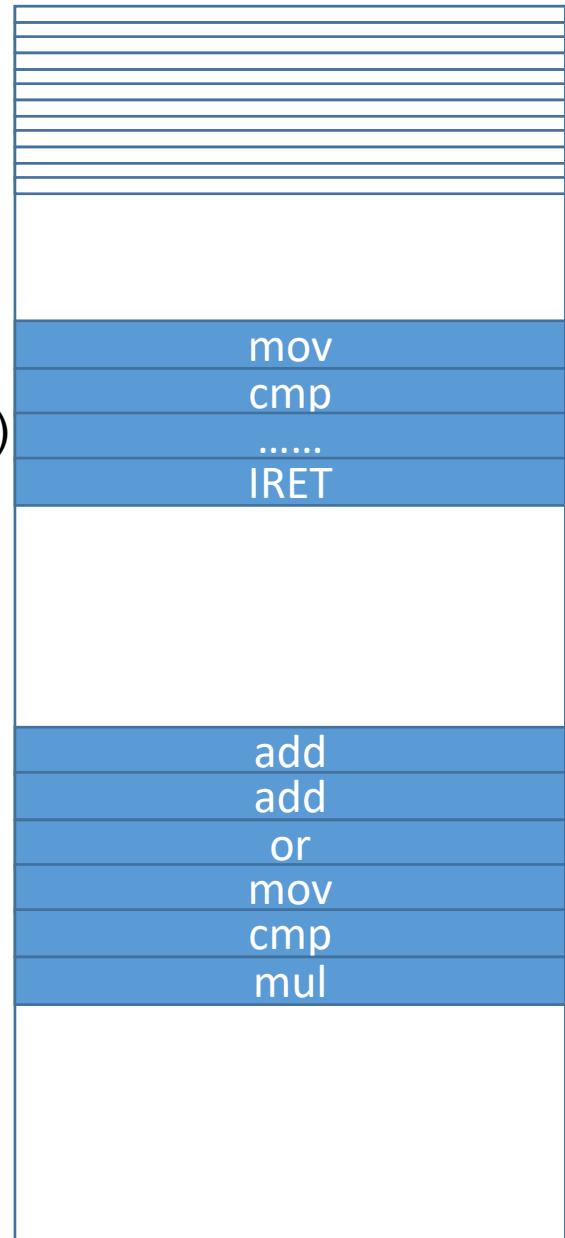
# Legacy PC Interrupt Mechanism



ISR  
(part of OS)

User  
program

IVT

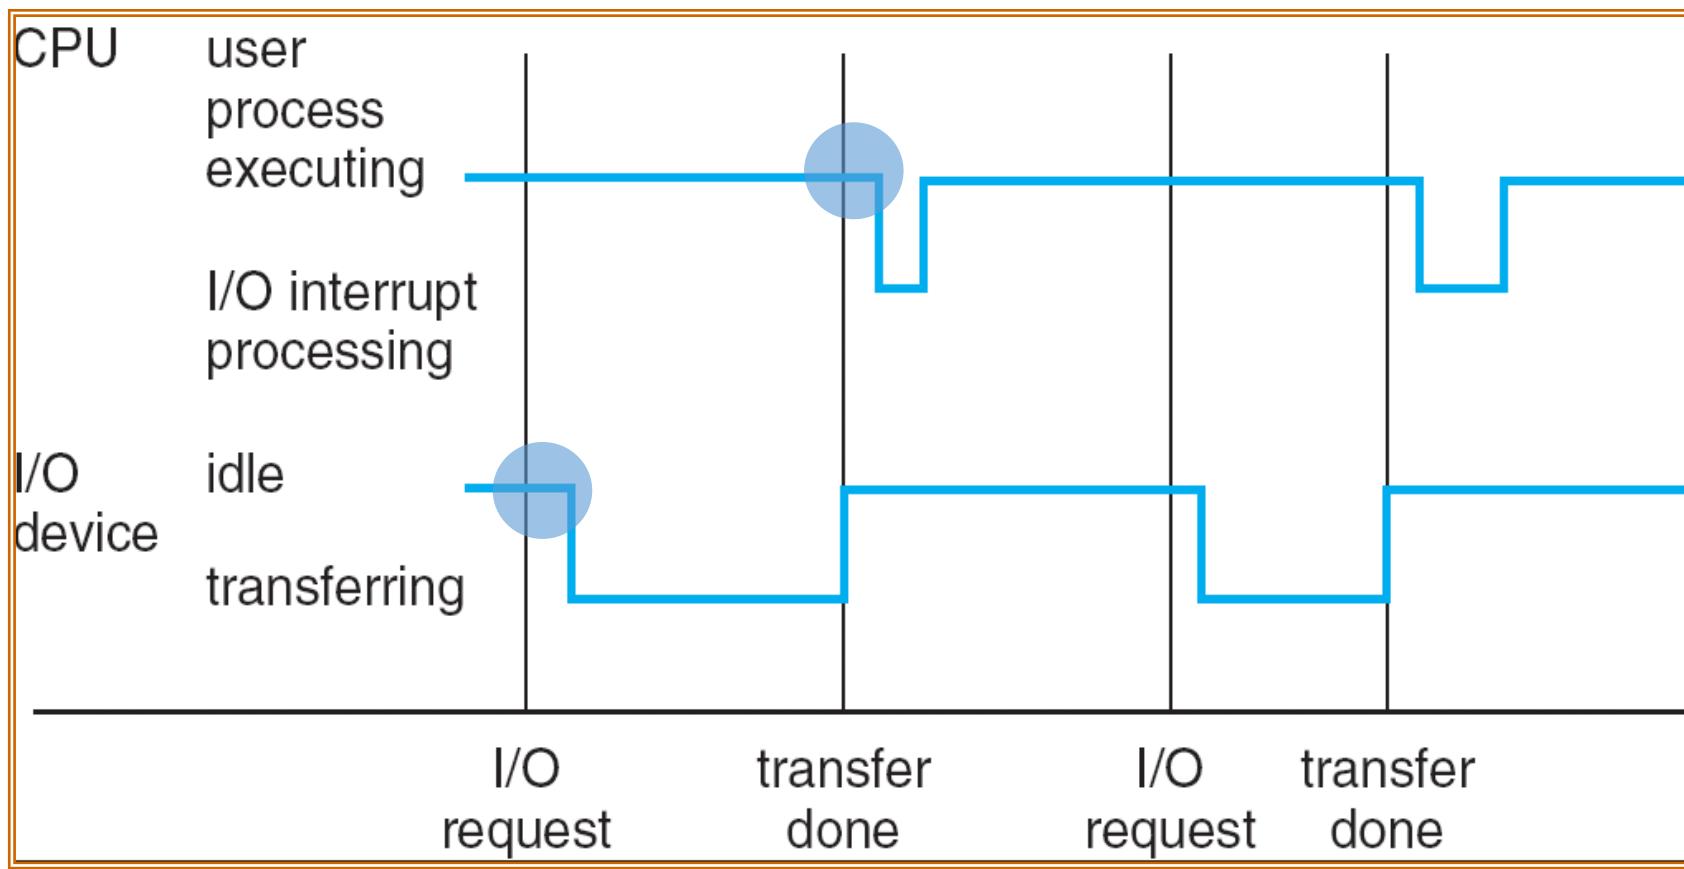


IVT=interrupt vector table

ISR=interrupt service routine

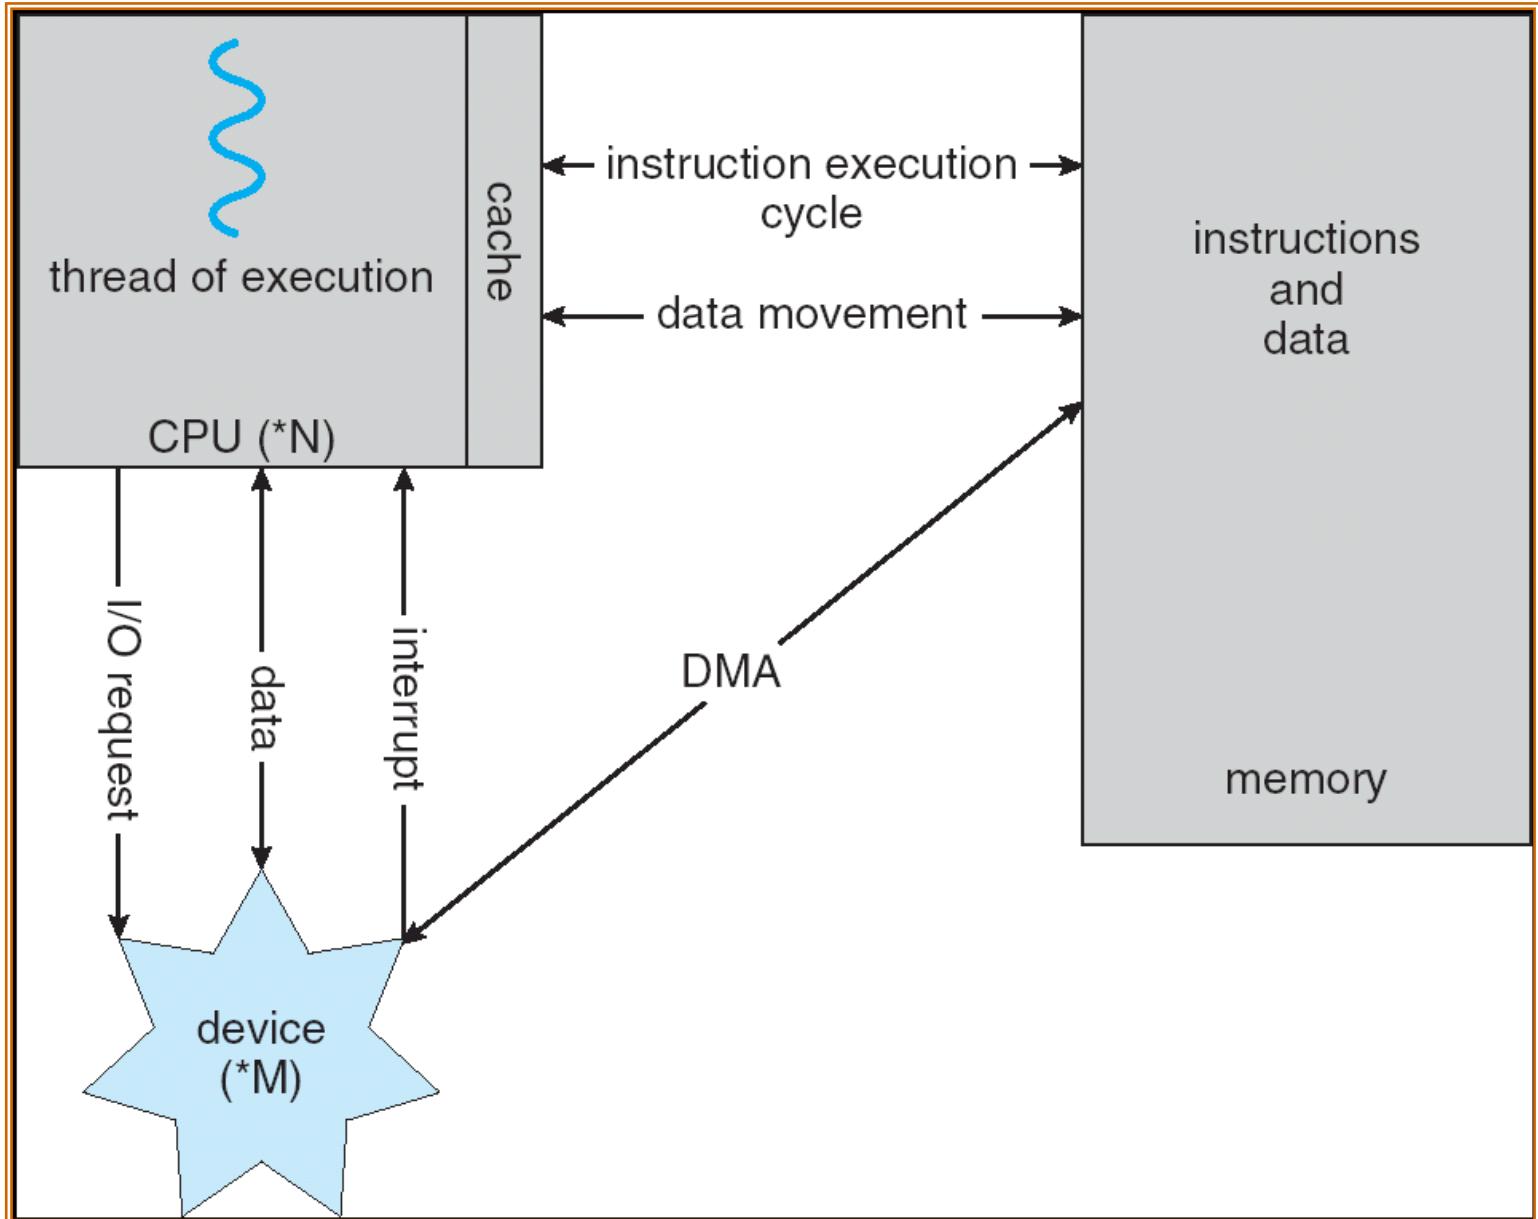
PIC=programmable interrupt controller

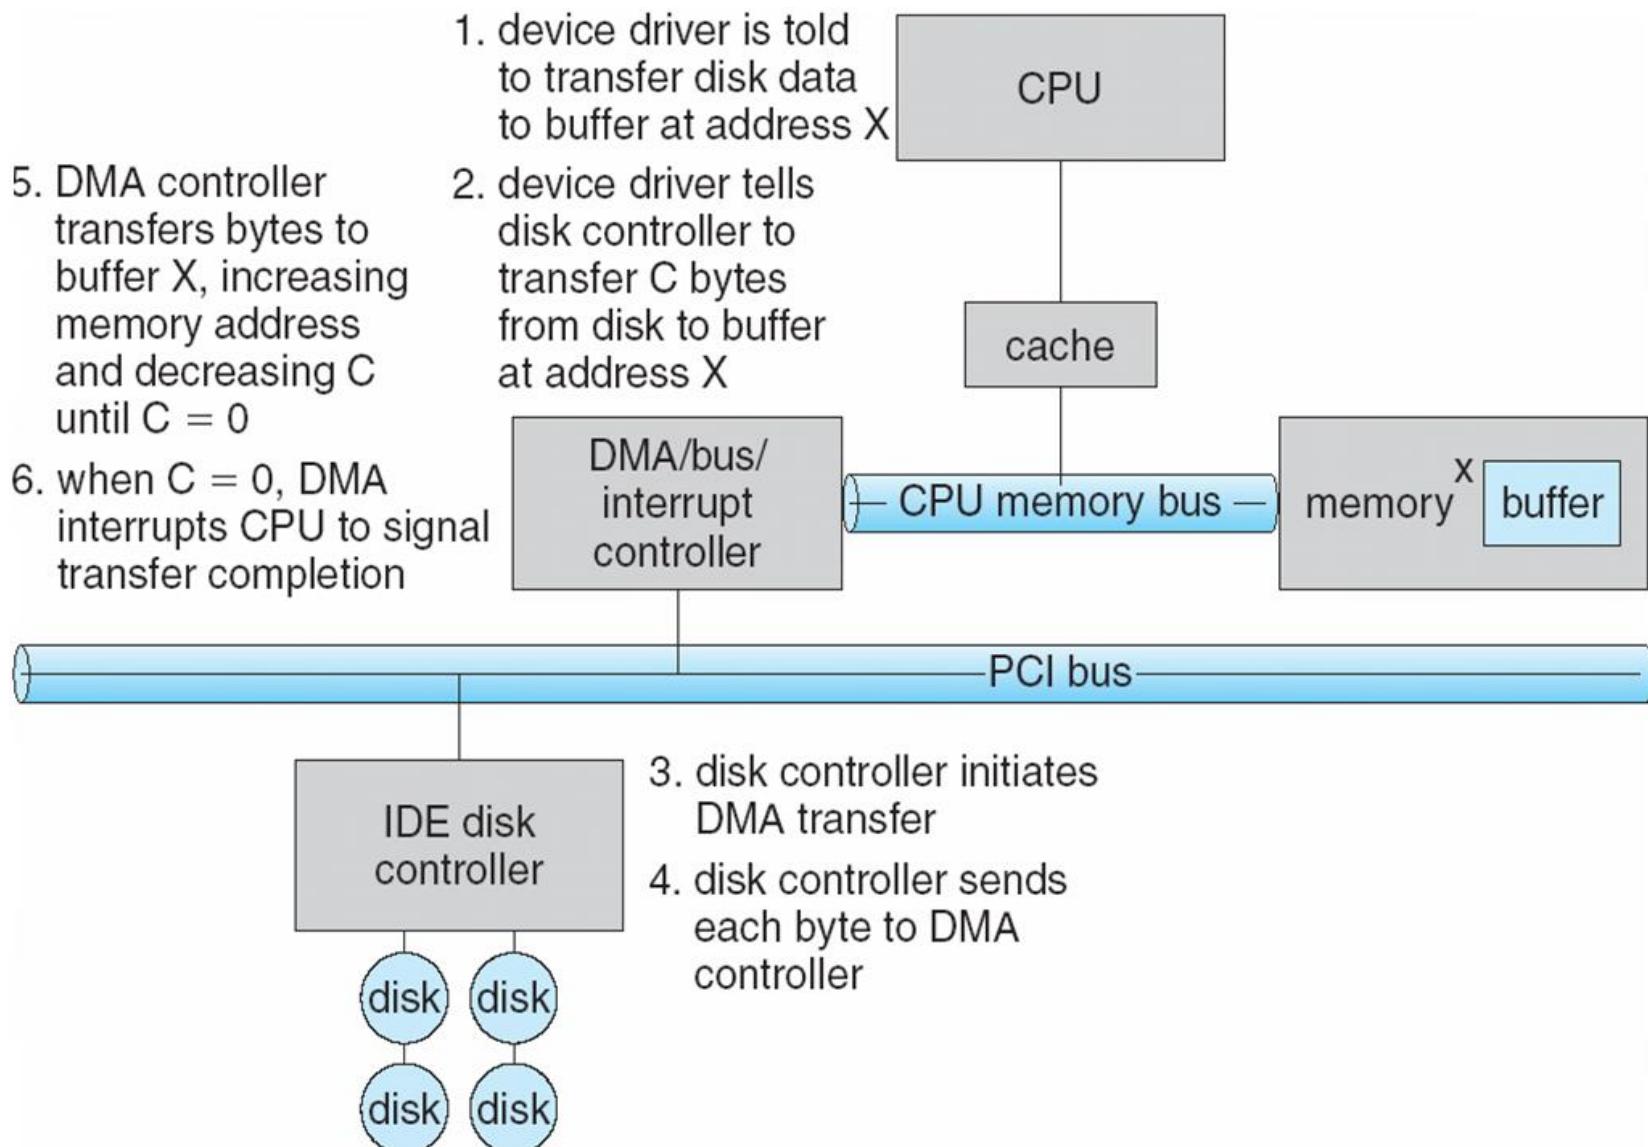
# Interrupt Timeline



# Direct Memory Access

- Data movement of I/O operation
  - Each device controller has a local buffer
  - Option 1: CPU moves data between main memory and I/O local buffers
  - Option 2: CPU offloads data movement to **DMA**
- DMA: Used for high-speed I/O devices able to transmit information at close to memory speeds
- Device controller transfers **blocks of data** from buffer storage directly to main memory without CPU intervention
- Only one interrupt is generated per block, rather than one interrupt per byte

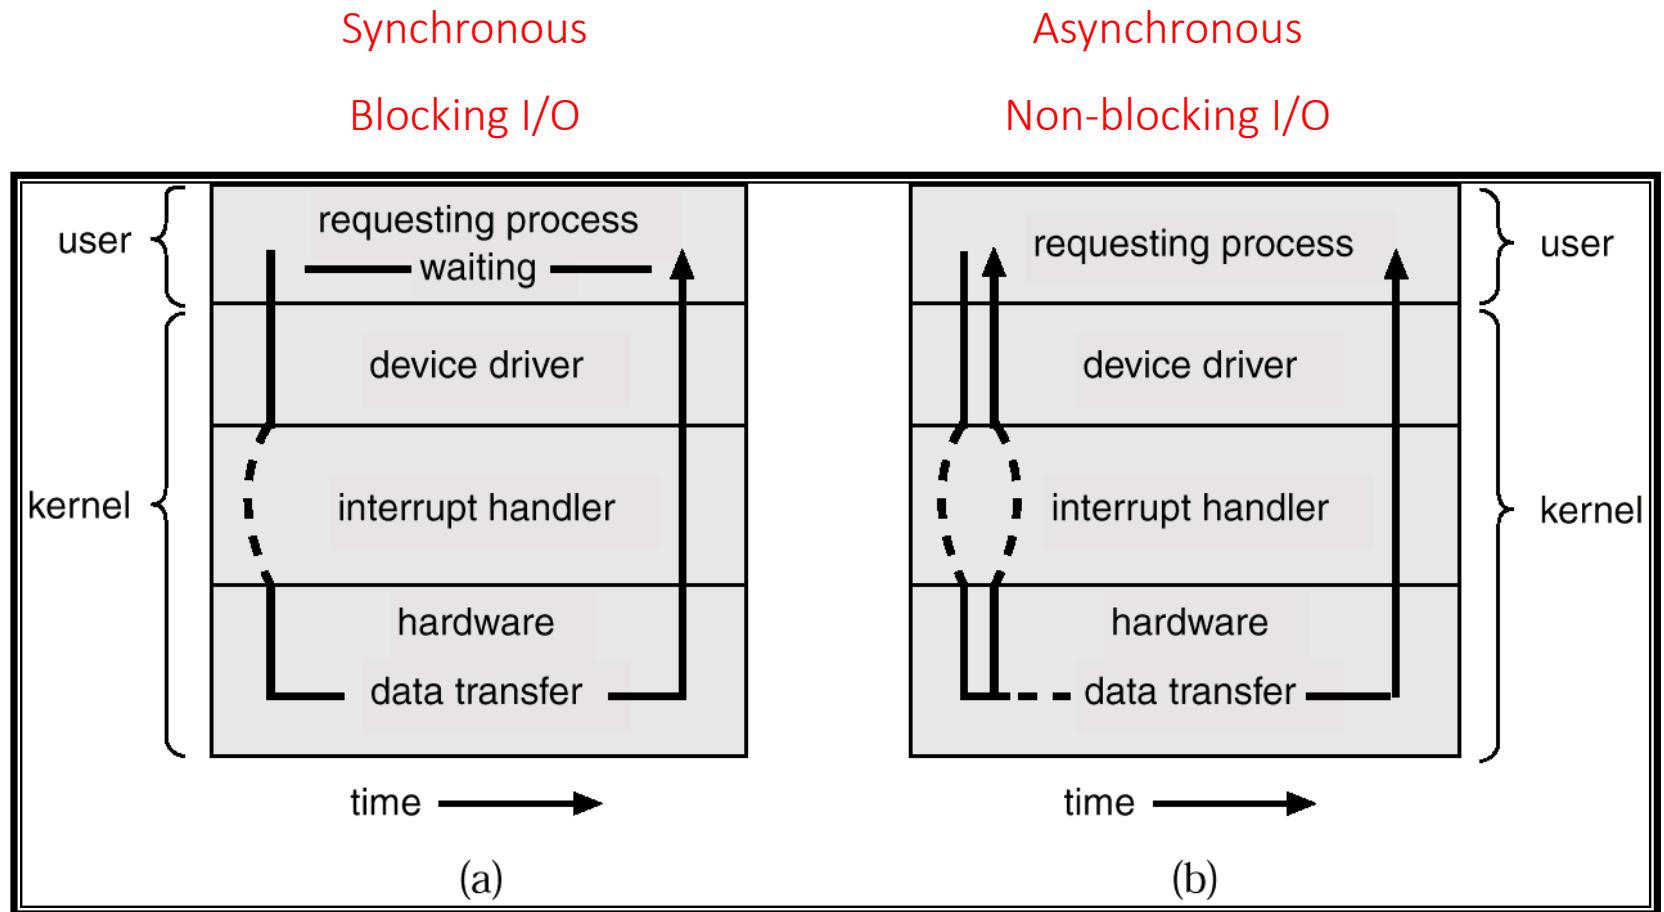




# Synchronous I/O vs. Asynchronous I/O

- After I/O starts, control returns to the program only upon I/O completion (**blocking call; sync I/O**)
  - Example: regular I/O reading `read()` or `fread()`
- After I/O starts, control returns to the program without waiting for I/O completion (**non-blocking call; async I/O**)
  - Example: regular I/O writing `write()` or `fwrite()`

# I/O Structure (Async vs Sync)

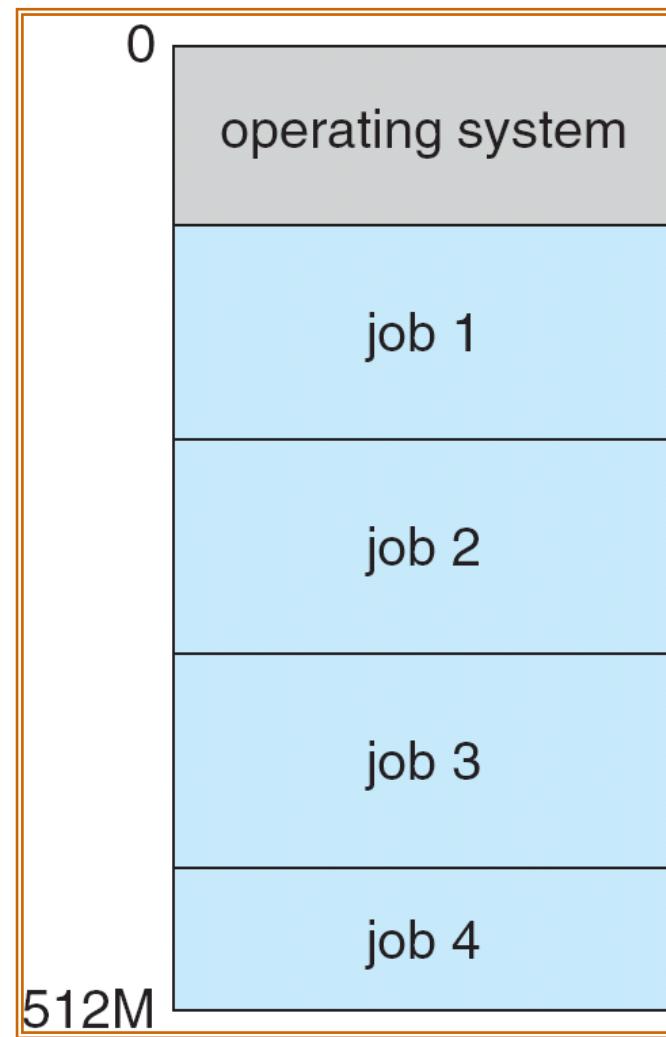


- Are the following operations synchronous or asynchronous?
  - `read()`
  - `write()`
  - `fsync()`
  - `aio_read()`

# Multiprogramming

and how it interacts with I/O operations

# A Possible Memory Layout of Multiprogramming Systems



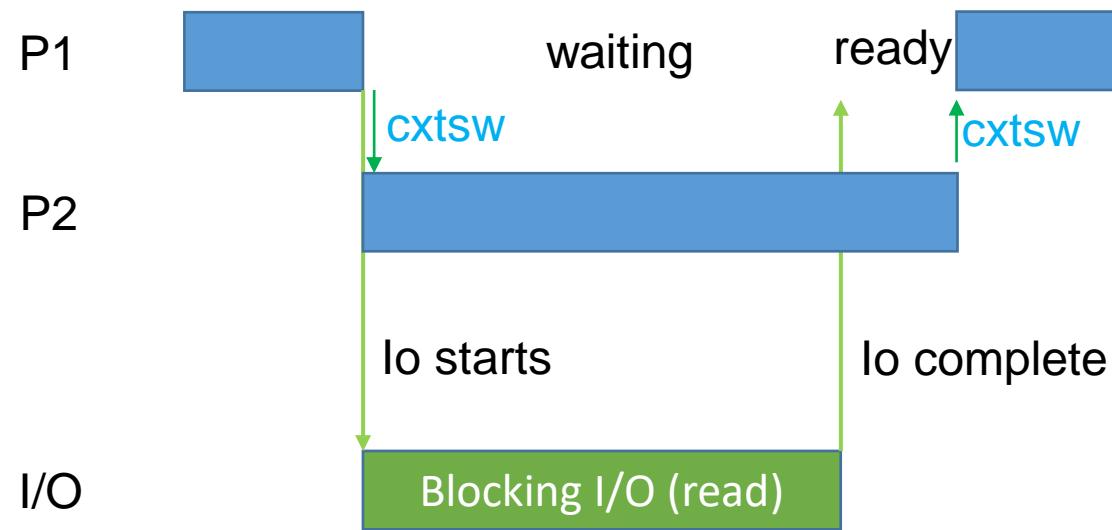
# Multiprogramming

- Single process **cannot keep CPU and I/O devices busy at all times**
- Multiprogramming organizes processes so CPU always has one to execute
- A **subset** of all processes in system is kept in memory
- One process selected and run via process scheduling
- When it has to wait (for I/O for example), OS switches to another process

# Timesharing

- Timesharing is logical extension in which **CPU switches processes so frequently** that users can interact with each job while it is running, creating **interactive computing**
  - Switching frequency usually  $\geq 100$  Hz, driven by timer IRQ
  - Each user has at least one program executing in memory  
⇒ process (ch3)
  - If several jobs ready to run at the same time ⇒ CPU scheduling (ch5)
  - If processes don't fit in memory, swapping moves them in and out to run (ch5)
  - Virtual memory allows execution of processes not completely in memory (ch8)

# Multitasking with I/Os



# Know the Terminology

- Multiprogramming
  - multiple processes are loaded into memory for execution
- Multitasking
  - Multiprogramming with overlapped task execution
- Timesharing
  - multitasking + periodic switch among processes

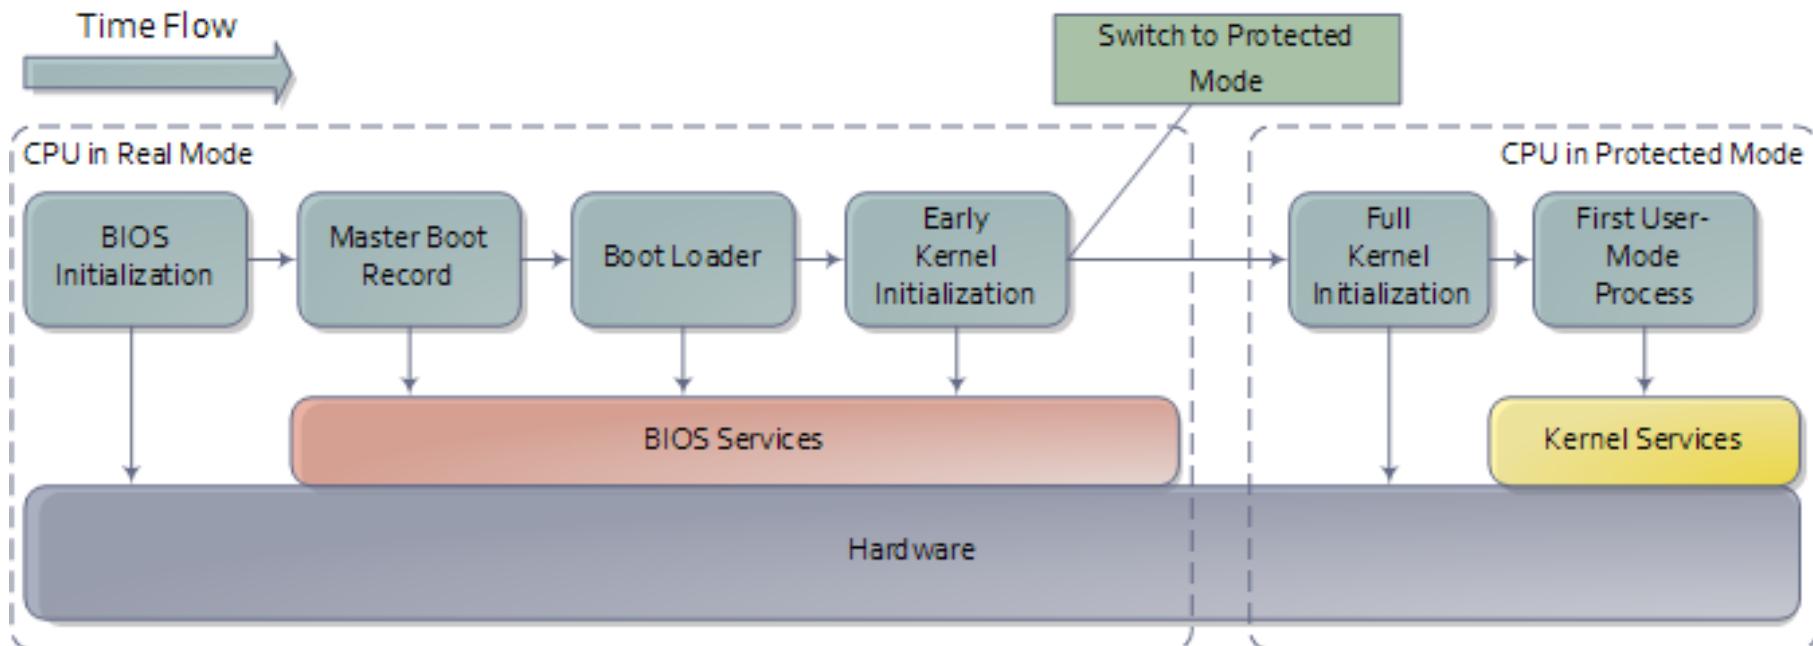
# Before the OS

# Before the OS: Computer Startup

- Bootstrap program is loaded at power-up or reboot
  - Typically stored in ROM or EEPROM, generally known as firmware
  - Initializes all aspects of system
  - Loads operating system kernel and starts execution
- Legacy PC bootloading procedure
  - FFFF:0000
  - BIOS
  - MBR (Master Boot Record)
  - Boot Manager
  - Pre-OS
  - OS

“BIOS+MBR” has been replaced by “UEFI+GPT”, an extended spec. However, their purposes are highly similar.

# PC Boot Sequence



Boot  
code

Absolute sector 0 (cylinder 0, head 0, sector 1)																
	0	1	2	3	4	5	6	7	8	9	A	B	C	D	E	F
0000	EB	48	90	00	00	00	00	00	00	00	00	00	00	00	00	00
0010	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00
0020	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00
0030	00	00	00	00	00	00	00	00	00	00	00	00	00	00	03	02
0040	80	00	00	80	DF	0A	93	01	00	08	FA	EA	50	7C	00	00
0050	31	CO	8E	D8	8E	D0	BC	00	20	FB	AO	40	7C	3C	FF	74
0060	02	88	C2	52	BE	76	7D	E8	34	01	F6	C2	80	74	54	B4
0070	41	BB	AA	55	CD	13	5A	52	72	49	81	FB	55	AA	75	43
0080	AO	41	7C	84	CO	75	05	83	E1	01	74	37	66	8B	4C	10
0090	BE	05	7C	C6	44	FF	01	66	8B	1E	44	7C	C7	04	10	00
00A0	C7	44	02	01	00	66	89	5C	08	C7	44	06	00	70	66	31
00B0	CO	89	44	04	66	89	44	0C	B4	42	CD	13	72	05	BB	00
00C0	70	EB	7D	B4	08	CD	13	73	0A	F6	C2	80	0F	84	F3	00
00D0	E9	8D	00	BE	05	7C	C6	44	FF	00	66	31	CO	88	F0	40
00E0	66	89	44	04	31	D2	88	CA	C1	E2	02	88	E8	88	F4	40
00F0	89	44	08	31	CO	88	D0	CO	E8	02	66	89	04	66	A1	44
0100	7C	66	31	D2	66	F7	34	88	54	0A	66	31	D2	66	F7	74
0110	04	88	54	0B	89	44	0C	3B	44	08	7D	3C	8A	54	0D	CO
0120	E2	06	8A	4C	0A	FE	C1	08	D1	8A	6C	0C	5A	8A	74	0B
0130	BB	00	70	8E	C3	31	DB	B8	01	02	CD	13	72	2A	8C	C3
0140	8E	06	48	7C	60	1E	B9	00	01	8E	DB	31	F6	31	FF	FC
0150	F3	A5	1F	61	FF	26	42	7C	BE	7C	7D	E8	40	00	EB	0E
0160	BE	81	7D	E8	38	00	EB	06	BE	8B	7D	E8	30	00	BE	90
0170	7D	E8	2A	00	EB	FE	47	52	55	42	20	00	47	65	6F	6D
0180	00	48	61	72	64	20	44	69	73	6B	00	52	65	61	64	00
0190	20	45	72	72	6F	72	00	BB	01	00	B4	0E	CD	10	AC	3C
01A0	00	75	F4	C3	00	00	00	00	00	00	00	00	00	00	00	00
01B0	00	00	00	00	00	00	00	00	A8	E1	A8	E1	00	00	80	01
01C0	01	00	07	FE	FF	6D	3F	00	00	00	AF	39	D7	00	00	00
01D0	C1	6E	0C	FE	FF	FF	EE	39	D7	00	BD	86	BB	00	00	FE
01E0	FF	FF	83	FE	FF	FF	AB	CO	92	01	CD	2F	03	00	00	FE
01F0	FF	FF	0F	FE	FF	FF	78	FO	95	01	83	AF	CC	00	55	AA

Partition  
table

Disk MBR that uses GRUB 0.92/0.93 boot code

MBR  
signature

# End of Chapter 1

# Review Questions

1. PCIe vs. PCI, which one is in the north bridge, and which one uses the serial protocol? Why?
2. Explain the entire process of interrupt handling, from the firing of a hardware event until the first line of code of the corresponding ISR is executed
3. Consider read() and write(). Whose call latency is subject to the speed of the I/O device? (hint: blocking and non-blocking)
4. Describe the boot procedure of the legacy PC