

Chapter 2: System Structures

Prof. Li-Pin Chang

CS@NYCU

Chapter 2: System Structures

- Operating System Services
 - User Interfaces
 - System Programs
 - **System Calls** and types of System Calls
- Operating System Design and Implementation
 - **Operating System Structure**
 - Monolithic kernels
 - Virtual Machines
 - Micro kernels

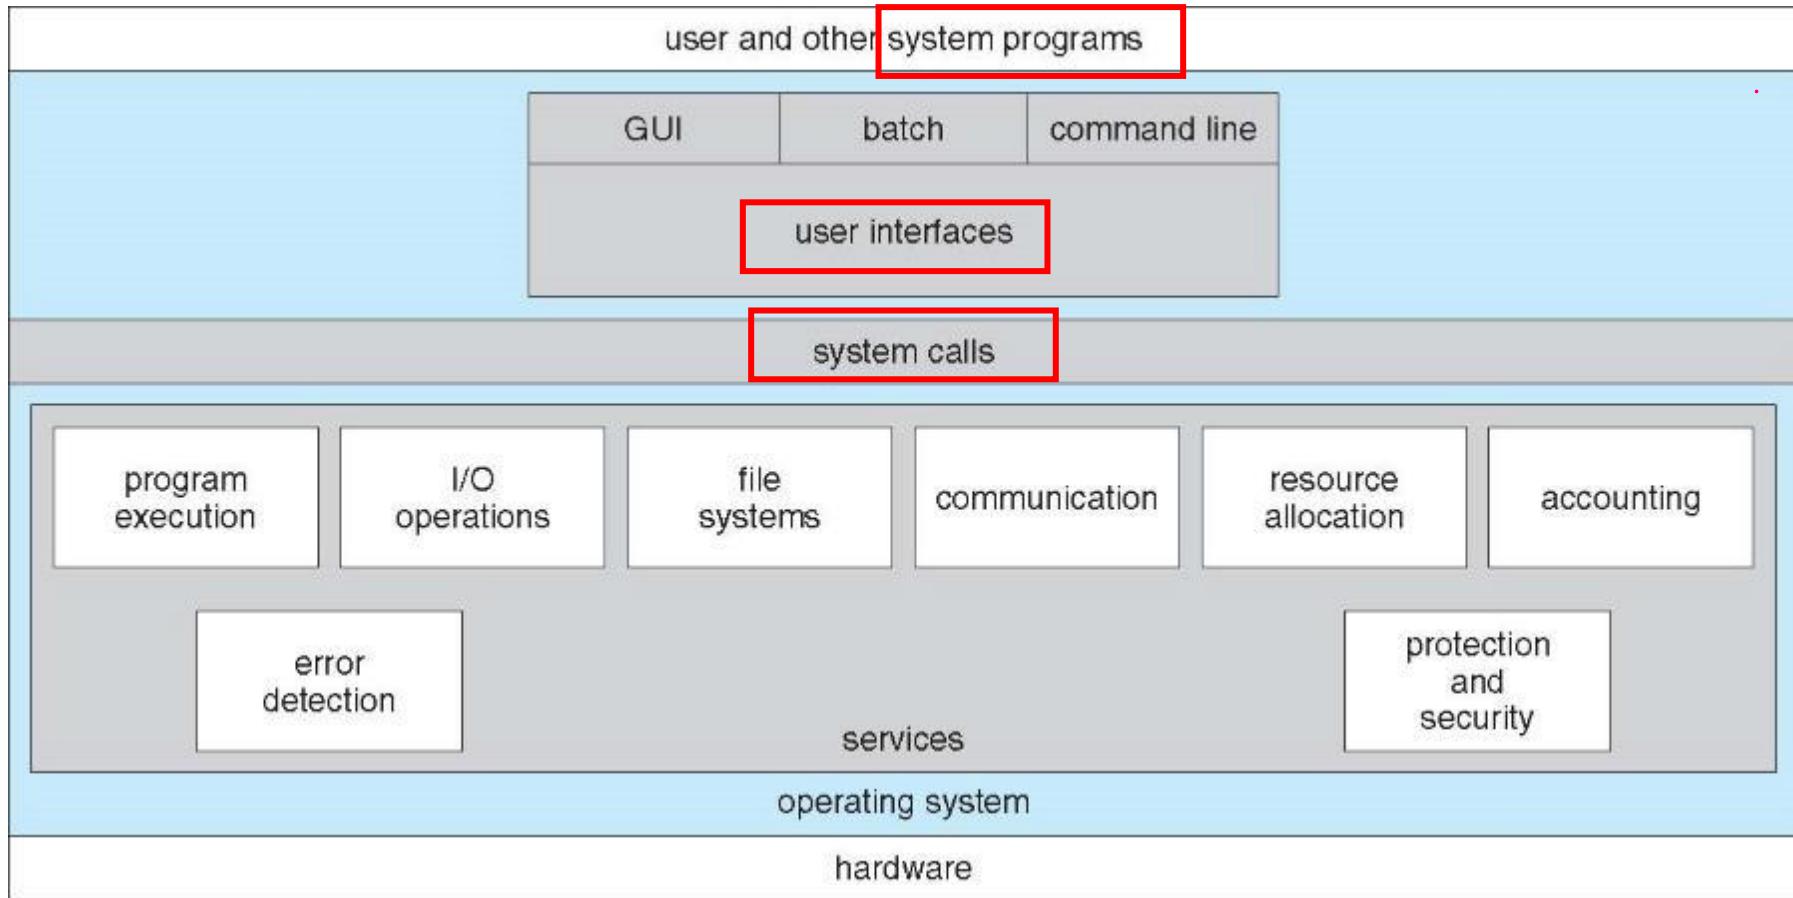
Objectives

- To describe the services an operating system provided to users, processes, and other systems
- To discuss the various ways of structuring an operating system

OPERATING SYSTEM SERVICES

A View of Operating System Services

3 levels of services



Operating System Services

- One set of operating-system services provides **functions** that are helpful to the user:
 - **User interface** - Almost all operating systems have a user interface (UI)
 - Varies between Command-Line (CLI), Graphics User Interface (GUI), Batch
 - **Program execution** - The system must be able to load a program into memory and to run that program, end execution, either normally or abnormally (indicating error)
 - **I/O operations** - A running program may require I/O, which may involve a file or an I/O device.
 - **File-system manipulation** - The file system is of particular interest. Obviously, programs need to read and write files and directories, create and delete them, search them, list file information, permission management.

Operating System Services (Cont.)

- One set of operating-system services provides **functions** that are helpful to the user (Cont):
 - **Communications** – Processes may exchange information, on the same computer or between computers over a network
 - Communications may be via shared memory or through message passing (packets moved by the OS)
 - **Error detection** – OS needs to be constantly aware of possible errors
 - May occur in the CPU and memory hardware, in I/O devices, in user program
 - For each type of error, OS should take the appropriate action to ensure correct and consistent computing
 - Debugging facilities can greatly enhance the user's and programmer's abilities to efficiently use the system

Operating System Services (Cont.)

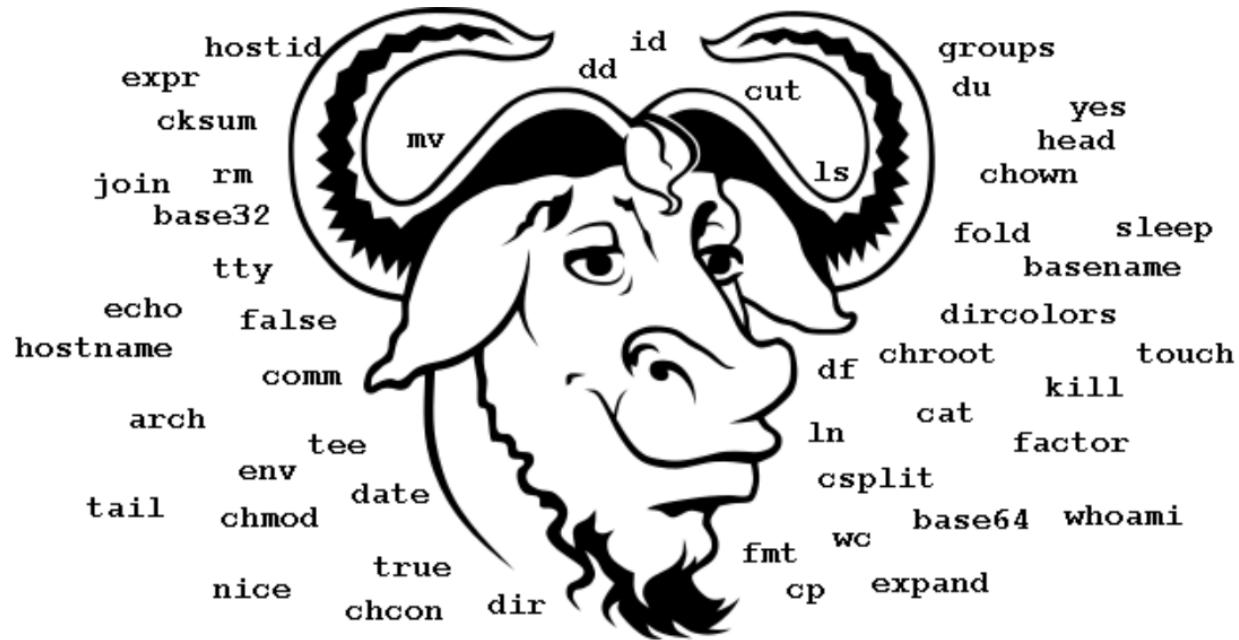
- Another set of OS functions exists for ensuring the **efficient operation** of the system itself via resource sharing
 - Resource **allocation** - When multiple users or multiple jobs running concurrently, resources must be allocated to each of them
 - Many types of resources - Some (such as CPU cycles, main memory, and file storage) may have special allocation code, others (such as I/O devices) may have general request and release code.
 - **Accounting** - To keep track of which users use how much and what kinds of computer resources
 - **Protection and security** - The owners of information stored in a multiuser or networked computer system may want to control use of that information, concurrent processes should not interfere with each other
 - Protection involves ensuring that all access to system resources is controlled
 - Security of the system from outsiders requires user authentication, extends to defending external I/O devices from invalid access attempts
 - If a system is to be protected and secure, precautions must be instituted throughout it. A chain is only as strong as its weakest link.

SYSTEM PROGRAMS

System Programs

- System programs provide a convenient environment for program development and execution. They can be divided into:
 - File manipulation (cp, mv...)
 - Status information (ls...)
 - File modification (vi...)
 - Programming language support (cc, as, ld, ar...)
 - Program loading and execution
 - Communications (telnet...)
- Most users' view of the operation system is defined by system programs, not the actual system calls

GNU coreutils + binutils



System Programs

- Provide a convenient environment for program development and execution
 - Some of them are simply user interfaces to system calls; others are considerably more complex
- File management - Create, delete, copy, rename, print, dump, list, and generally manipulate files and directories
- Programming-language support - **Compilers, assemblers, debuggers** and interpreters sometimes provided
- Program loading and execution- Absolute loaders, relocatable loaders, linkage editors, and overlay-loaders, debugging systems for higher-level and machine language

USER INTERFACE

Operating-System User Interface - CLI

- Shell refers to the interface program between users and the kernel
 - Text-driven: Command Line Interface (CLI)
 - Graphics-driven: Graphical User Interface (GUI)
- CLI allows direct command entry
 - Primarily fetches a command from user and executes it
 - Sometimes commands built-in, sometimes just names of programs
 - If the latter, adding new features doesn't require shell modification

CLI in Windows/Linux

```
C:\>cmd
C:\windows\system32>cd \

C:\>dir/w
磁碟區 C 中的磁碟是 TI30940600B
磁碟區序號: 0A90-10B7

C:\ 的目錄

[BaKoMa Tex] [EcpaComponent]
[Intell] [LJP1100_P1560_P1600_Full_Solution]
[PerfLogs] [Program Files]
[Program Files (x86)] SSUUpdater.log
[TOSHIBA] [Users]
[UTDService] [Windows]
               1 個檔案      282 位元組
               11 個目錄  55,639,674,880 位元組可用

C:\>_
```

```
Loading...

Welcome to JS/Linux (x86)

Use 'vflogin username' to connect to your account.
You can create a new account at https://vfsync.org/signup .
Use 'export_file filename' to export a file to your computer.
Imported files are written to the home directory.

[root@localhost ~]# ls -l
total 8
drwxr-xr-x    3 root      root          163 Aug 21  2011 dos
-rw-r--r--    1 root      root         242 Jul 15  2017 hello.c
[root@localhost ~]# pwd
/root
[root@localhost ~]# █
```

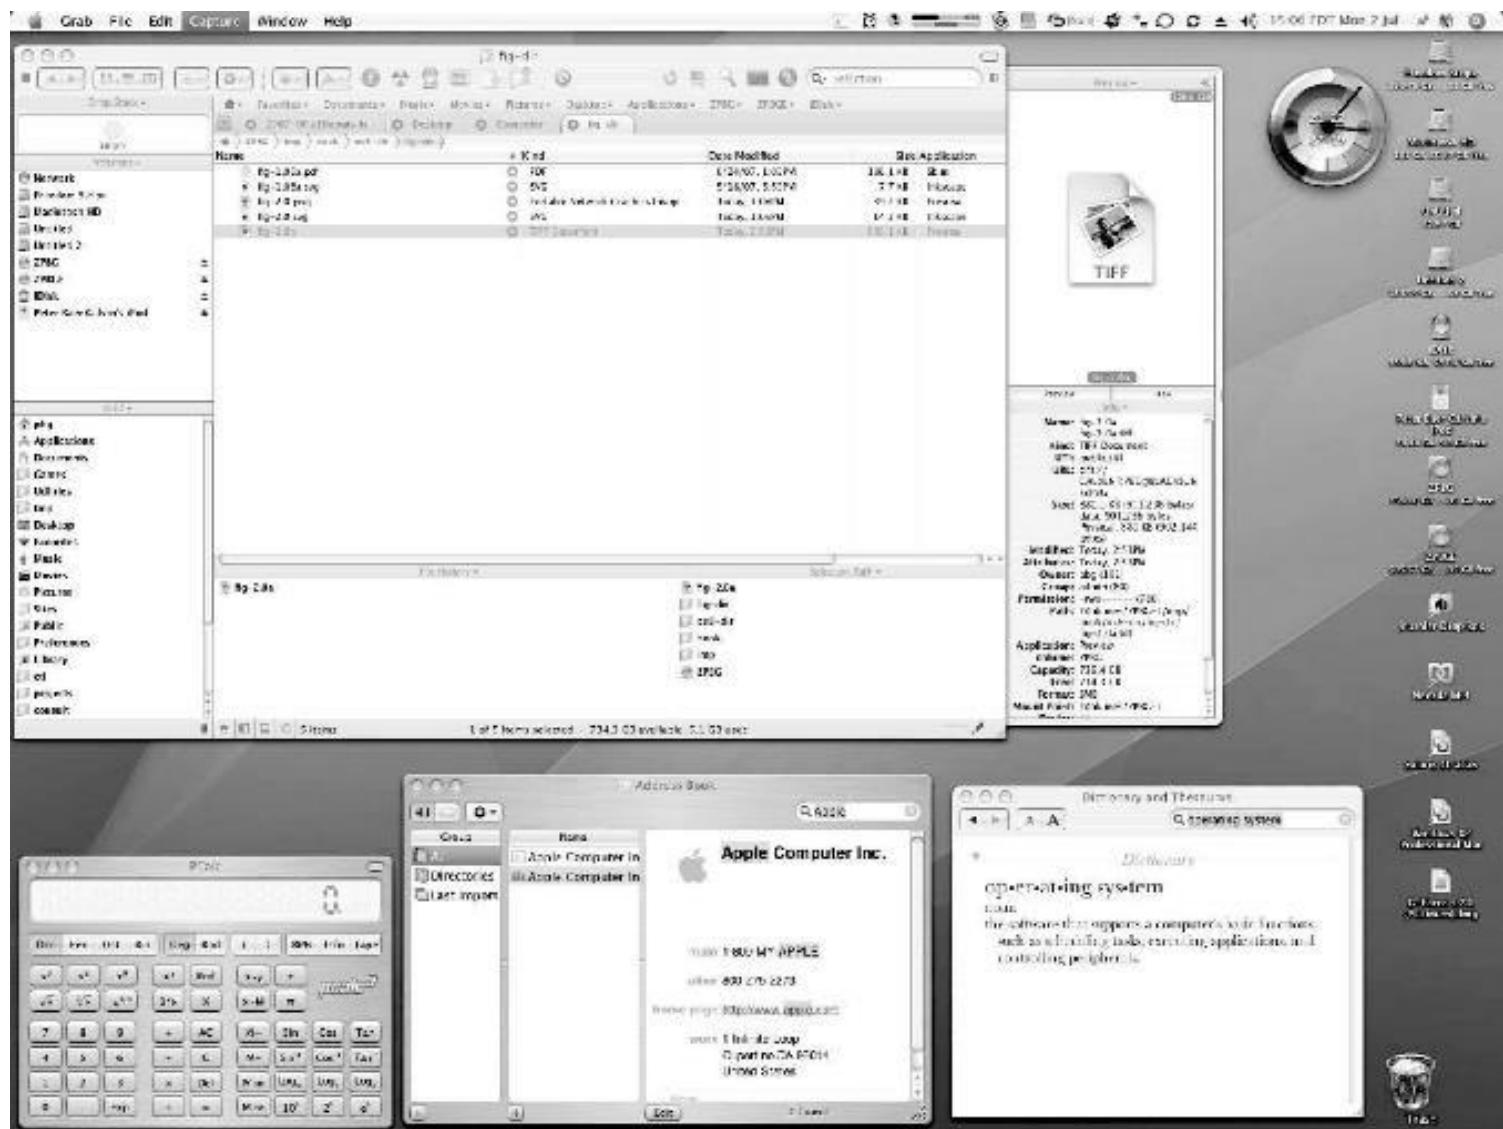
User Operating System Interface - GUI

- User-friendly desktop metaphor interface
 - Usually mouse, keyboard, and monitor
 - Icons represent files, programs, actions, etc
 - Invented by Xerox PARC
 - Sarcasm: Window, Icon, Menu, Pointer -- WIMP
- Many systems now include both CLI and GUI interfaces
 - Microsoft Windows is GUI with CLI “command” shell
 - Apple Mac OS X as “Aqua” GUI interface with UNIX kernel underneath and shells available
 - Solaris is CLI with optional GUI interfaces (Java Desktop, KDE)

Xerox PARC's Achievements

- 1971 – Laser printer
- 1973 – Alto personal computer (GUI, mouse, Ethernet)
- 1973 – Ethernet (local area networking)
- 1973 – Smalltalk (object-oriented programming)
- 1975 – Practical mouse for GUI
- 1975 – Bravo, first WYSIWYG editor
- 1980s – InterPress (precursor to PostScript)
PDF is an encapsulation of postscript

The Mac OS X GUI



Touchscreen Interfaces

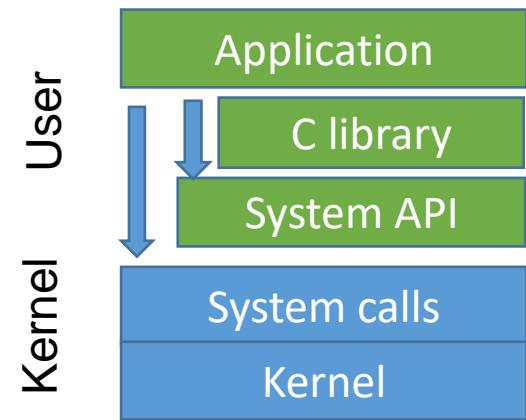
- Touchscreen devices require new interfaces
 - Mouse not possible or not desired
 - Actions and selection based on **gestures**
 - Virtual keyboard for text entry



SYSTEM CALLS

System Calls

- Programming interface to the services provided by the OS
- Mostly accessed by programs via a system Application Program Interface (**API**) rather than direct system call use
 - Portability and simplicity
- Three most common system APIs are **Win32 API** for Windows, **POSIX API** for UNIX, and **Java API** for the Java VM



Standard C library:

`fopen("w+..."`)

C language



WIN32 API:

`CreateFile()`

Windows >= win4.0, >=95



Kernel API:

`NTCreateFile()`

WinNT, 2k,XP, vista



System Call:

`int 2e`

X86 machine instruction

Std C lib (stdlib)

```
int printf ( const char * format, ... );
```



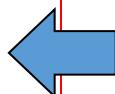
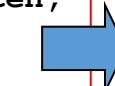
```
BOOL WINAPI WriteFile( Win32 API  
    _In_             HANDLE hFile,  
    _In_             LPCVOID lpBuffer,  
    _In_             DWORD nNumberOfBytesToWrite,  
    _Out_opt_        LPDWORD lpNumberOfBytesWritten,  
    _Inout_opt_      LPOVERLAPPED lpOverlapped  
);
```

```
fread (ucrtbase)  
-> _fread_nolock/_read_nolock  
-> ReadFile (KERNEL32/KERNELBASE)  
-> NtReadFile (ntdll, user-mode)  
-> syscall  
-> NtReadFile/ZwReadFile (ntoskrnl, kernel-mode)  
-> IRP_MJ_READ → 檔案系統→儲存控制器→裝置
```

```
NTSTATUS NtWriteFile ( Kernel API  
    HANDLE          hFile,  
    HANDLE          hEvent,  
    PIO_APC_ROUTINE apc,  
    void*           apc_user,  
    PIO_STATUS_BLOCK io_status,  
    const void*     buffer,  
    ULONG           length,  
    PLARGE_INTEGER  offset,  
    PULONG          key  
)
```

System call (x86 CPU)

```
mov eax, <service #>  
lea edx, <addr of 1st arg>  
int 2e
```



Example of System API (Win32)

- Consider the ReadFile() function in the Win32 API—a function for reading from a file

The diagram shows the C language signature for the ReadFile() function. It is enclosed in a rectangular box with a double orange border. Inside, the return value is labeled "return value" with a downward arrow. The function name "ReadFile" is at the top, followed by its parameters: "c (HANDLE file, LPVOID buffer, DWORD bytes To Read, LPDWORD bytes Read, LPOVERLAPPED ovl);". A bracket on the right side of the parameters is labeled "parameters". An upward arrow points from the word "function name" to the start of the function signature.

```
return value
      ↓
• BOOL ReadFile c (HANDLE file,
                   LPVOID buffer,
                   DWORD bytes To Read,
                   LPDWORD bytes Read,
                   LPOVERLAPPED ovl);
                   ↑
function name
```

- A description of the parameters passed to ReadFile()
 - HANDLE file—the file to be read
 - LPVOID buffer—a buffer where the data will be read into and written from
 - DWORD bytesToRead—the number of bytes to be read into the buffer
 - LPDWORD bytesRead—the number of bytes read during the last read
 - LPOVERLAPPED ovl—indicates if overlapped I/O is being used

Example of System API (POSIX)

EXAMPLE OF STANDARD API

As an example of a standard API, consider the `read()` function that is available in UNIX and Linux systems. The API for this function is obtained from the `man` page by invoking the command

```
man read
```

on the command line. A description of this API appears below:

```
#include <unistd.h>

ssize_t      read(int fd, void *buf, size_t count)
```

return function parameters
value name

A program that uses the `read()` function must include the `unistd.h` header file, as this file defines the `ssize_t` and `size_t` data types (among other things). The parameters passed to `read()` are as follows:

- `int fd`—the file descriptor to be read
- `void *buf`—a buffer where the data will be read into
- `size_t count`—the maximum number of bytes to be read into the buffer

On a successful read, the number of bytes read is returned. A return value of 0 indicates end of file. If an error occurs, `read()` returns `-1`.

System Call Implementation

- Typically, **a number** associated with each system call
 - System-call interface maintains **a table** indexed according to these numbers
- The system call interface invokes intended system call in OS kernel and returns status of the system call and any return values
- The caller need know nothing about how the system call is implemented
 - Just needs to obey API and understand what OS will do as a result call

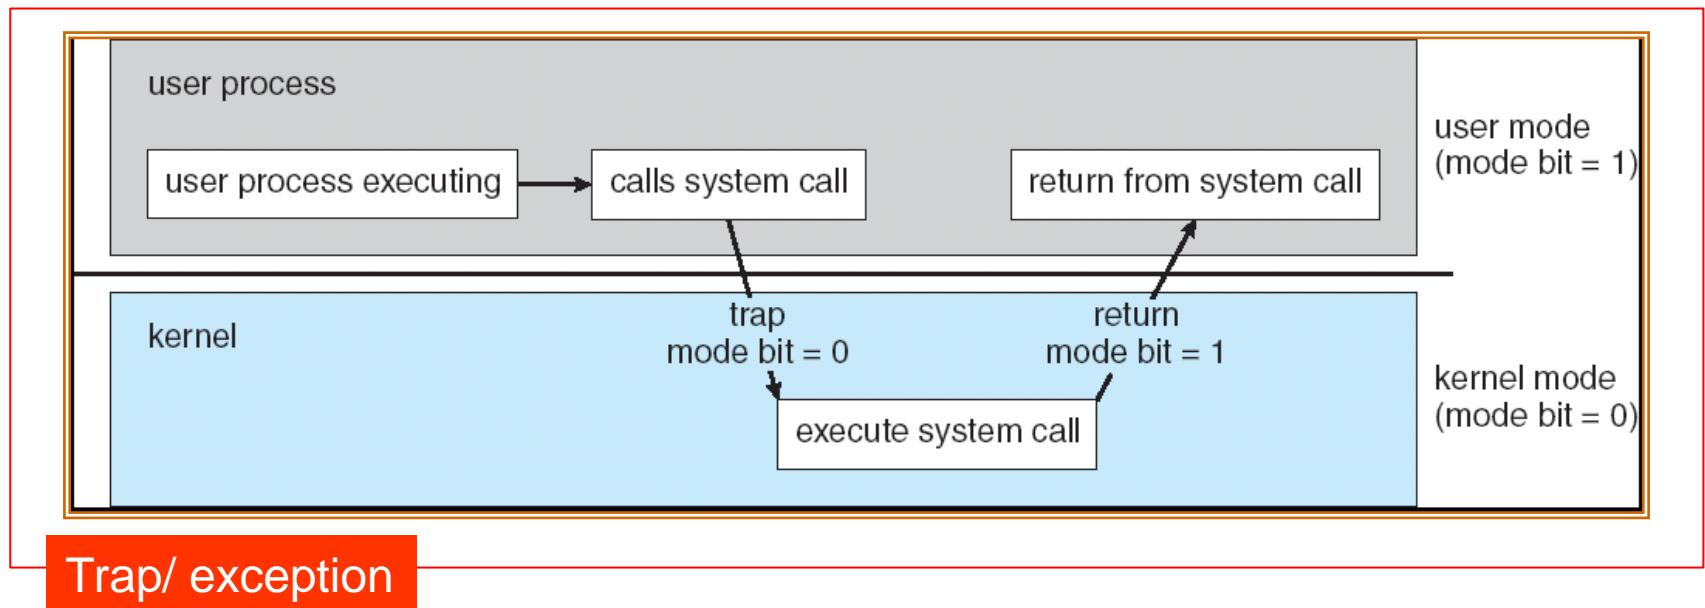
Dual Mode Operations

- Application calls into the kernel through software interrupt
 - also known as **trap** or **exception**
 - Software error: Div by zero, memory access violation, etc
 - Request for operating system service (system calls)

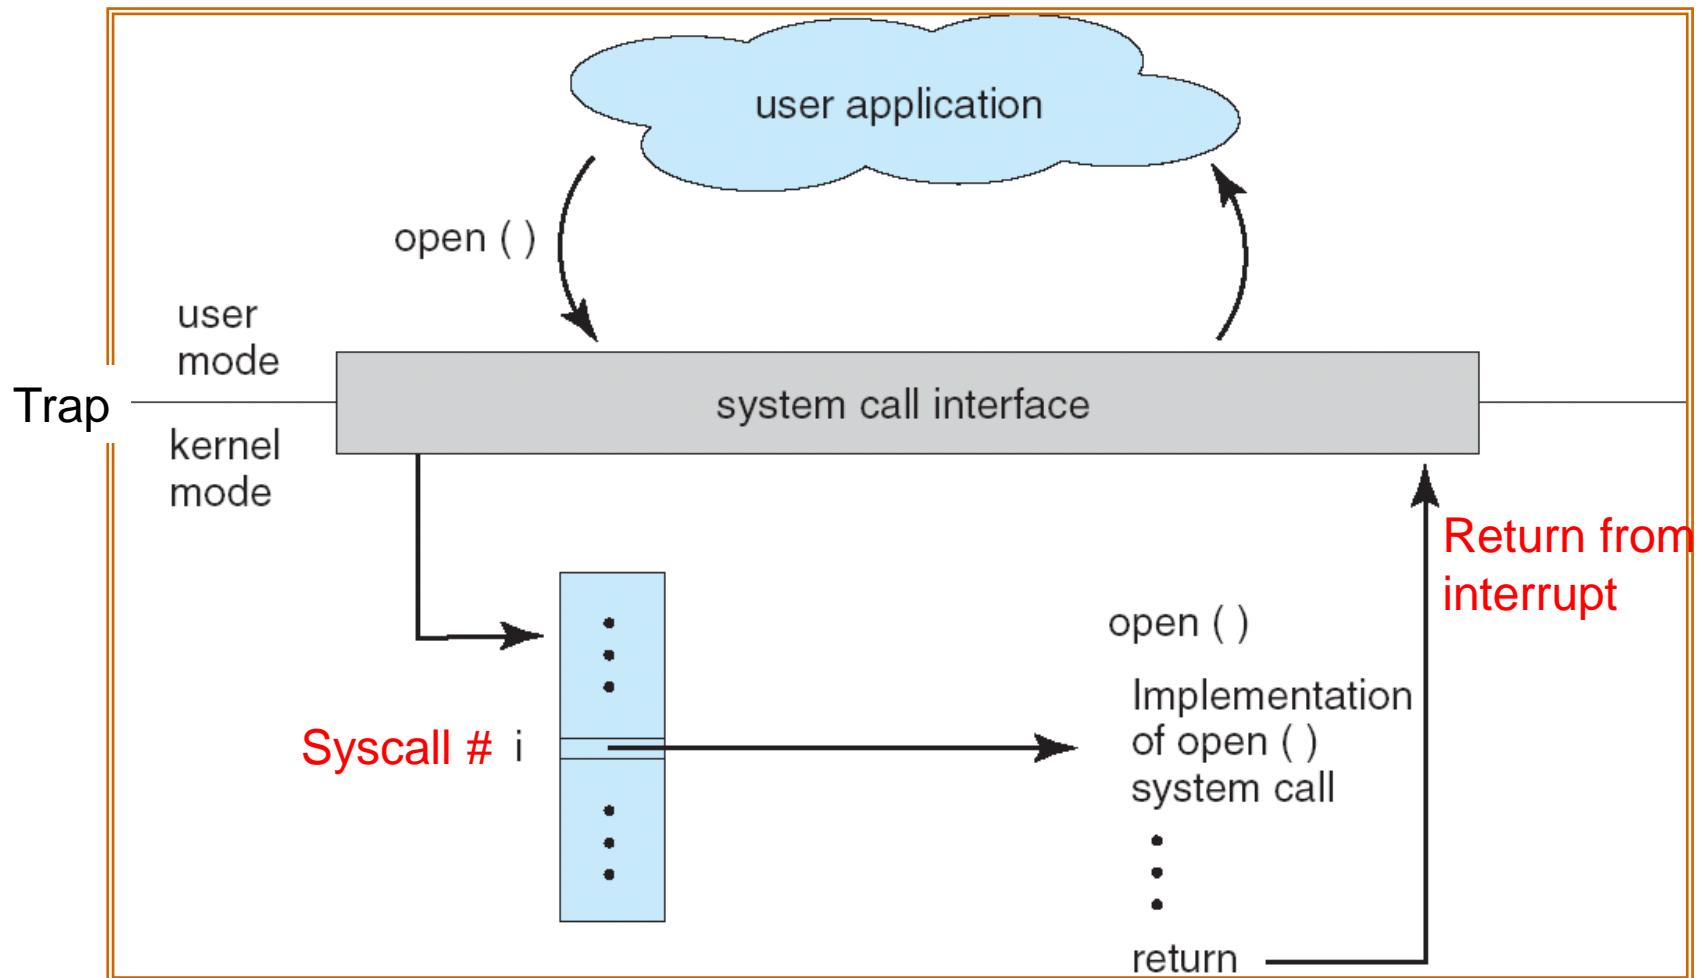
Dual Mode Operations

- Dual-mode operation, supported by CPU *hardware*, allows OS to protect itself against un-trusted code
 - User mode (untrusted, limited privilege)
 - Kernel mode (trusted, full privilege)
- The purpose of dual-mode design
 - Provides ability to distinguish when system is running user code or kernel code
 - Some instructions designated as *privileged*, only executable in kernel mode
 - System call changes mode to kernel, return from call resets it to user (mode bit supported by the CPU)

Transition from User to Kernel Mode



API – System Call – OS Relationship



System Call Parameter Passing

- A system call accepts an operation code and a set of parameters
- Three general methods to pass parameters to the OS
 - Pass the parameters in **registers**
 - Parameters stored in a block, or **table**, in memory, and address of block passed as a parameter in a register
 - Parameters placed, or pushed, into the **stack** by the program and popped off the stack by the operating system

Direct System Call in Linux (NASM Syntax)

```
section .data          ;declare section
msg db  "Hello World! :)",0xa ;our dear string
len equ $ - msg        ;length of our dear string

section .text          ; section declaration

    global _start      ; exporting entry point
                      ; to the ELF linker

_start:
; write Hello World string
    mov edx,len ;third arg: message length
    mov ecx,msg ;second arg: pointer to message to write
    mov ebx,1   ;first arg: file handle (stdout)
    mov eax,4   ;system call nr. (sys_write)
    int 0x80    ;call kernel (trigger a trap)

; and exit
    mov ebx,0   ;first syscall args: exit code
    mov eax,1   ;system call no. (sys_exit)
    int 0x80    ;call kernel
```

More on System Calls

- Direct System Call in Windows NT
 - Use the following fragment of assembly code to call the kernel

```
MOV EAX, <service #>
LEA EDX, <addr of 1st arg>
INT 2E
```

- Return value is in EAX (if any)
- For modern Intel / AMD CPUs, SYSENTER / SYSCALL are suggested for making system calls, respectively
 - Skip IVT lookup (dest. addr stored in a control register)
 - Fewer register backup

TYPES OF SYSTEM CALLS

Types of System Calls

- Process control
- File management
- Device management
- Information maintenance
- Communications

Examples of Windows and Unix System Calls

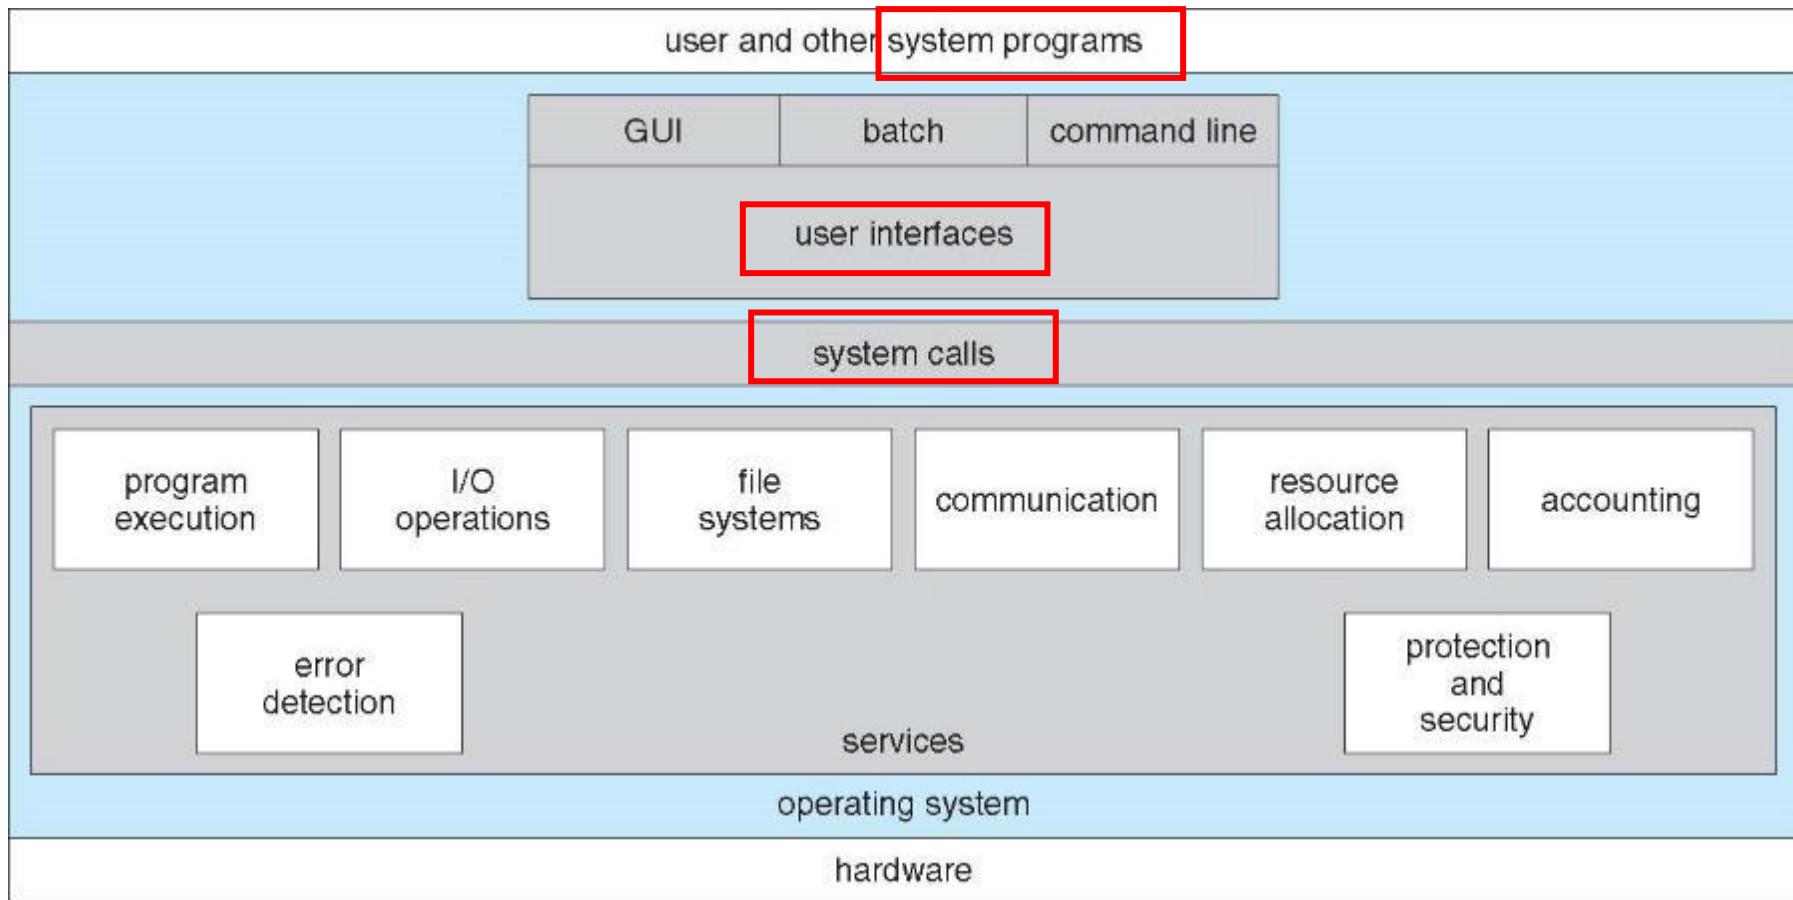
	Windows	Unix
Process Control	CreateProcess() ExitProcess() WaitForSingleObject()	fork() exit() wait()
File Manipulation	CreateFile() ReadFile() WriteFile() CloseHandle()	open() read() write() close()
Device Manipulation	SetConsoleMode() ReadConsole() WriteConsole()	ioctl() read() write()
Information Maintenance	GetCurrentProcessID() SetTimer() Sleep()	getpid() alarm() sleep()
Communication	CreatePipe() CreateFileMapping() MapViewOfFile()	pipe() shmget() mmap()
Protection	SetFileSecurity() InitializeSecurityDescriptor() SetSecurityDescriptorGroup()	chmod() umask() chown()

Linux System Calls

List by system call number

00 sys_setup [sys_ni_syscall]	70 sys_setreuid	140 sys_llseek [sys_lseek]
01 sys_exit	71 sys_setregid	141 sys_getdents
02 sys_fork	72 sys_sigsuspend	142 sys_newselect [sys_select]
03 sys_read	73 sys_sigpending	143 sys_flock
04 sys_write	74 sys_sethostname	144 sys_msync
05 sys_open	75 sys_setrlimit	145 sys_readyv
06 sys_close	76 sys_getrlimit	146 sys_writev
07 sys_waitpid	77 sys_getrusage	147 sys_getsid
08 sys_creat	78 sys_gettimeofday	148 sys_fdatasync
09 sys_link	79 sys_settimeofday	149 sys_sysctl [sys_sysctl]
10 sys_unlink	80 sys_getgroups	150 sys_mlock
11 sys_execve	81 sys_setgroups	151 sys_munlock
12 sys_chdir	82 sys_select [old_select]	152 sys_mlockall
13 sys_time	83 sys_symlink	153 sys_munlockall
14 sys_mknod	84 sys_oldlstat [sys_lstat]	154 sys_sched_setparam
15 sys_chmod	85 sys_readlink	155 sys_sched_getparam
16 sys_lchown	86 sys_uselib	156 sys_sched_setscheduler
17 sys_break [sys_ni_syscall]	87 sys_swapon	157 sys_sched_getscheduler
18 sys_oldstat [sys_stat]	88 sys_reboot	158 sys_sched_yield
19 sys_lseek	89 sys_readdir [old_readdir]	159 sys_sched_get_priority_max
20 sys_getpid	90 sys_mmap [old_mmap]	160 sys_sched_get_priority_min
21 sys_mount	91 sys_munmap	161 sys_sched_rr_get_interval

Recap: Operating System Services



OPERATING SYSTEM DESIGN AND IMPLEMENTATION

Operating System Design and Implementation

- Design and Implementation of OS not “solvable”, but some approaches have proven successful
- Internal structure of different Operating Systems can vary widely
- Start by defining **goals** and **specifications**
- Affected by choice of hardware, type of system
- User goals and System goals
 - User goals –convenient to use, easy to learn, reliable, safe, and fast
 - System goals –easy to design, implement, and maintain, as well as flexible, reliable, error-free, and efficient
- Design issues for different types of systems
 - Real-time OS: predictable latency, romability (small footprint)
 - Mainframe OS: throughput, scalability
 - Desktop OS: interaction, user friendly

Operating System Design and Implementation (Cont.)

- Important principle to separate
 - Mechanism: How to do it?
 - Policy: What will be done?
- Mechanisms determine how to do something, policies decide what will be done next
- The separation of policy from mechanism is a very important principle, it allows maximum flexibility if policy decisions are to be changed later
- Use disk I/O as an example:
 - Mechanism: How to read and write from disk?
 - Policy: Which disk I/O operation should be performed first?

Which one(s) of the following are policies; which are mechanisms?

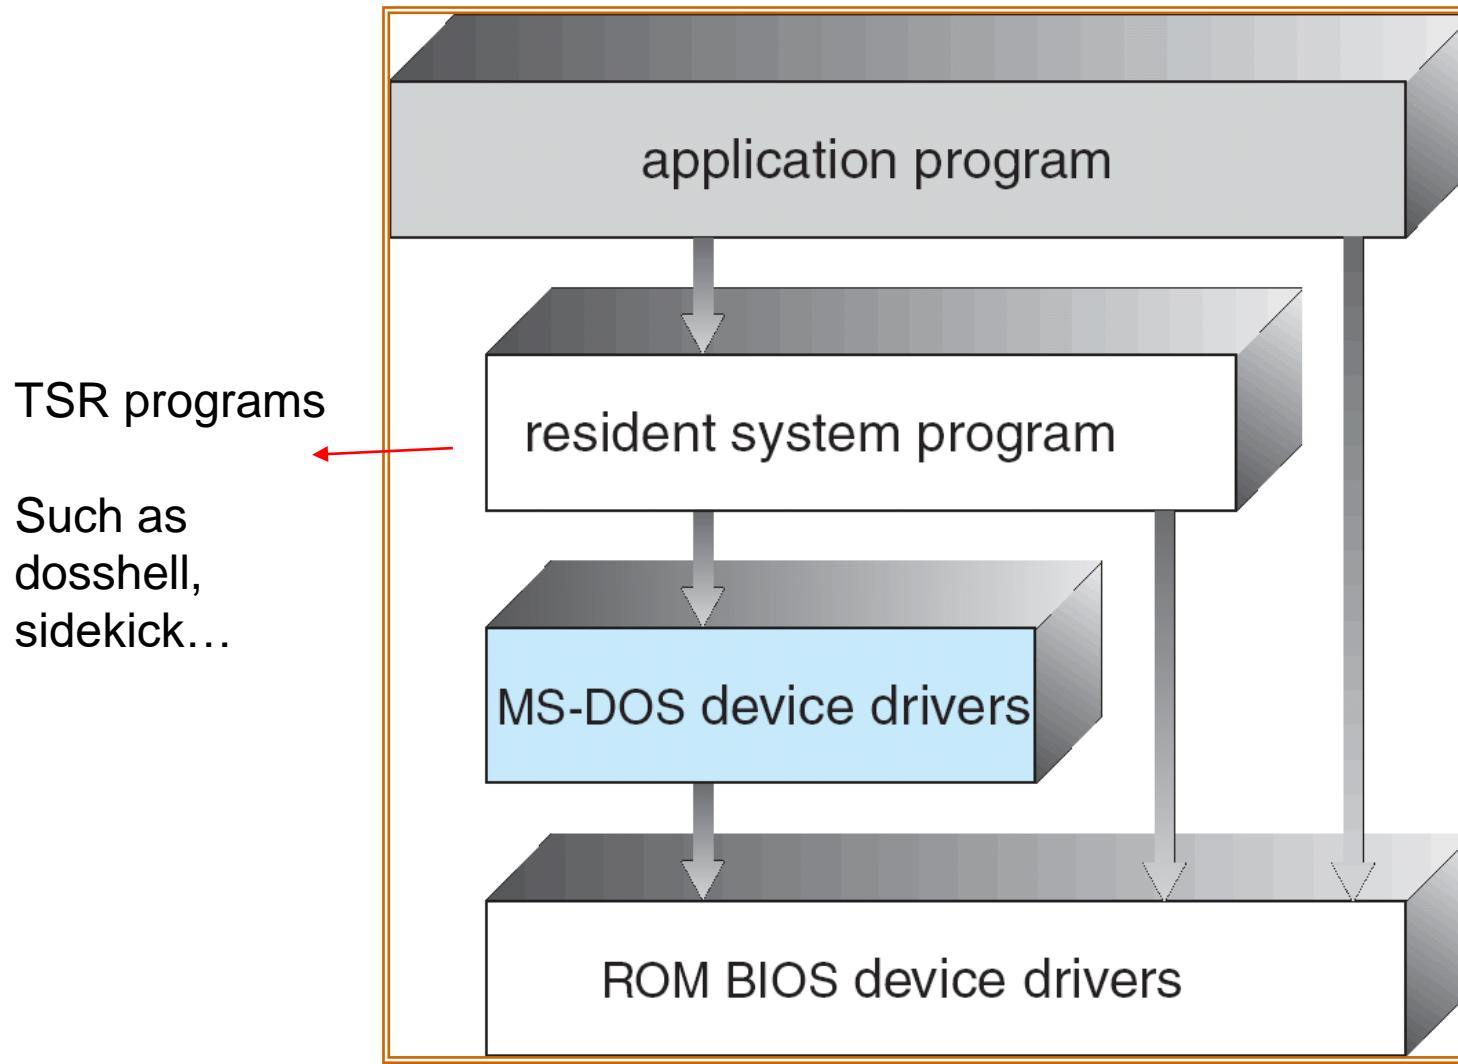
- a) process suspend/resume
- b) allocating the smallest among the memory blocks which are larger than the requested size
- c) marking a disk block as allocated
- d) servicing the disk I/O request which is closest to the disk head

OPERATING-SYSTEM STRUCTURE

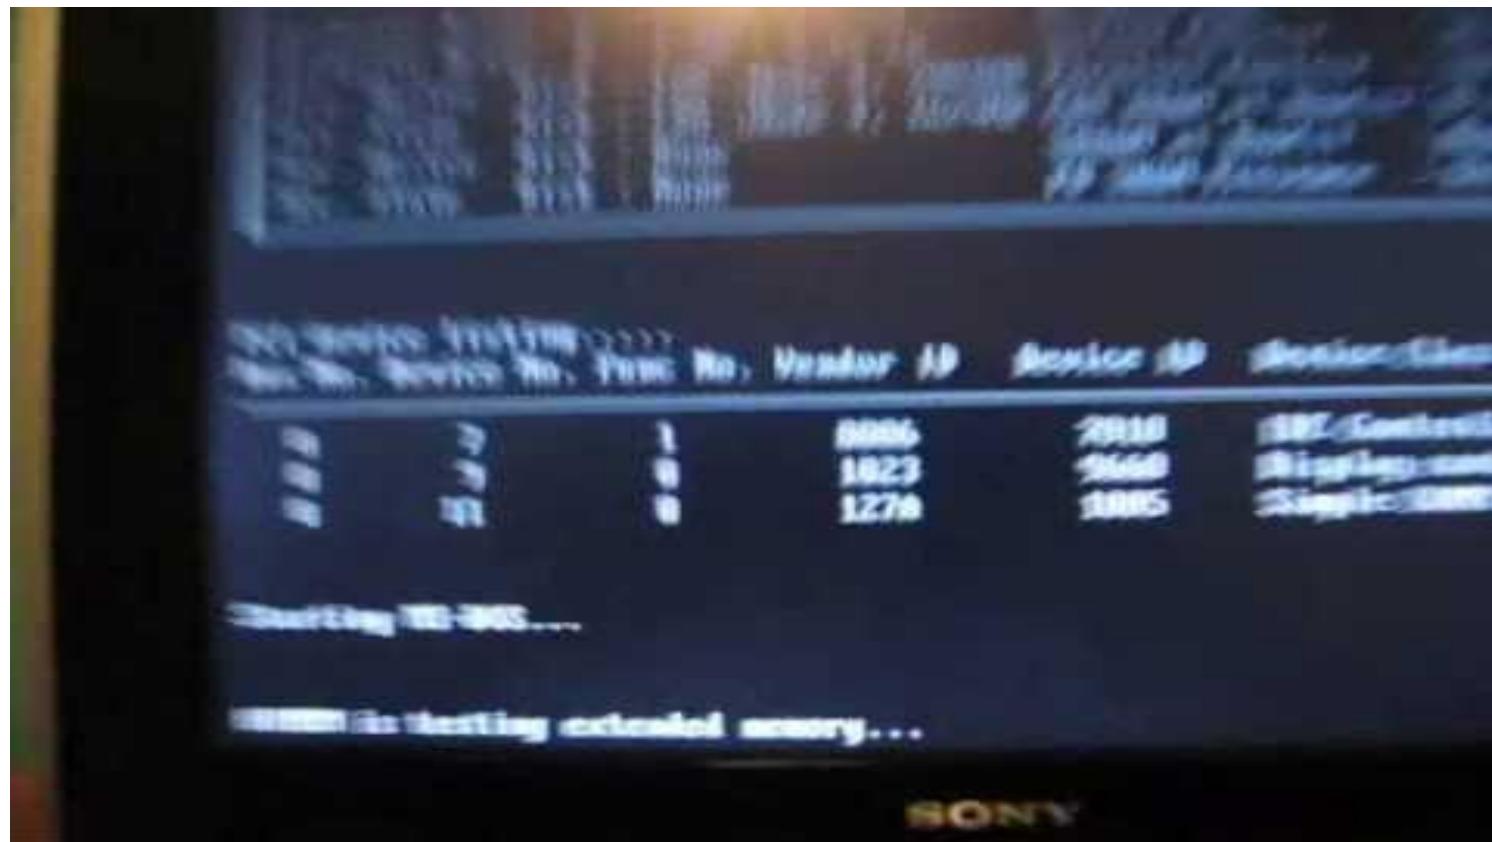
Simple Structure

- MS-DOS – written to provide the most functionality in the least space
 - Although MS-DOS has some structure, its interfaces and levels of functionality are not well separated
- No protection. MS-DOS is designed for 8086, which is not capable of any protection (e.g., user kernel)
 - Applications can directly access any memory addresses and control hardware (dangerous)
 - Near bare-metal performance (cool)
- Still alive and kicking, e.g., FreeDOS used in simple tasks such as motherboard firmware update

MS-DOS Layer Structure

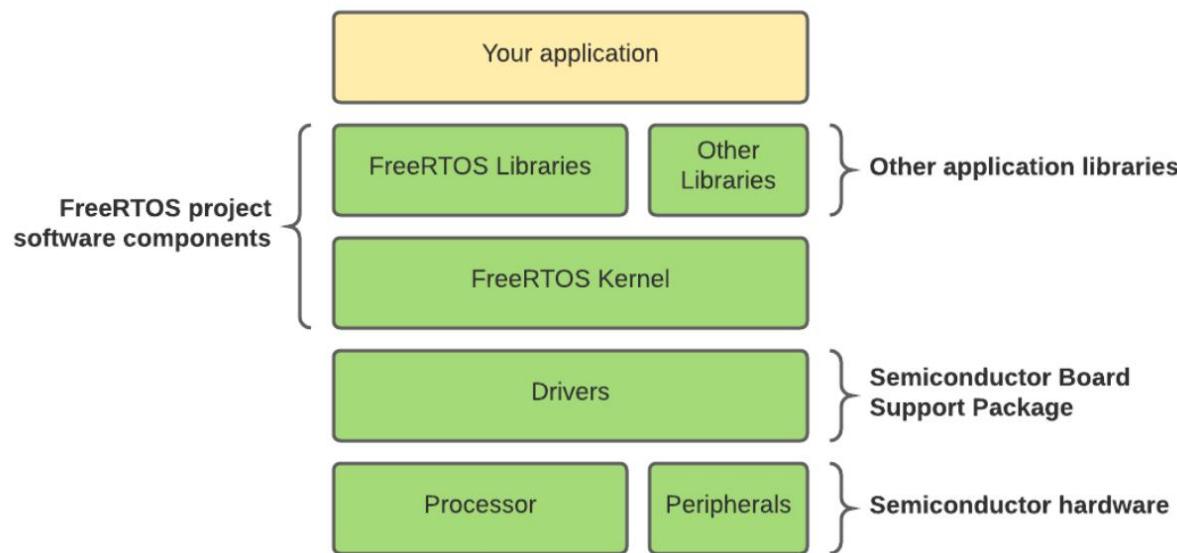


Booting MS-DOS



FreeRTOS Structure

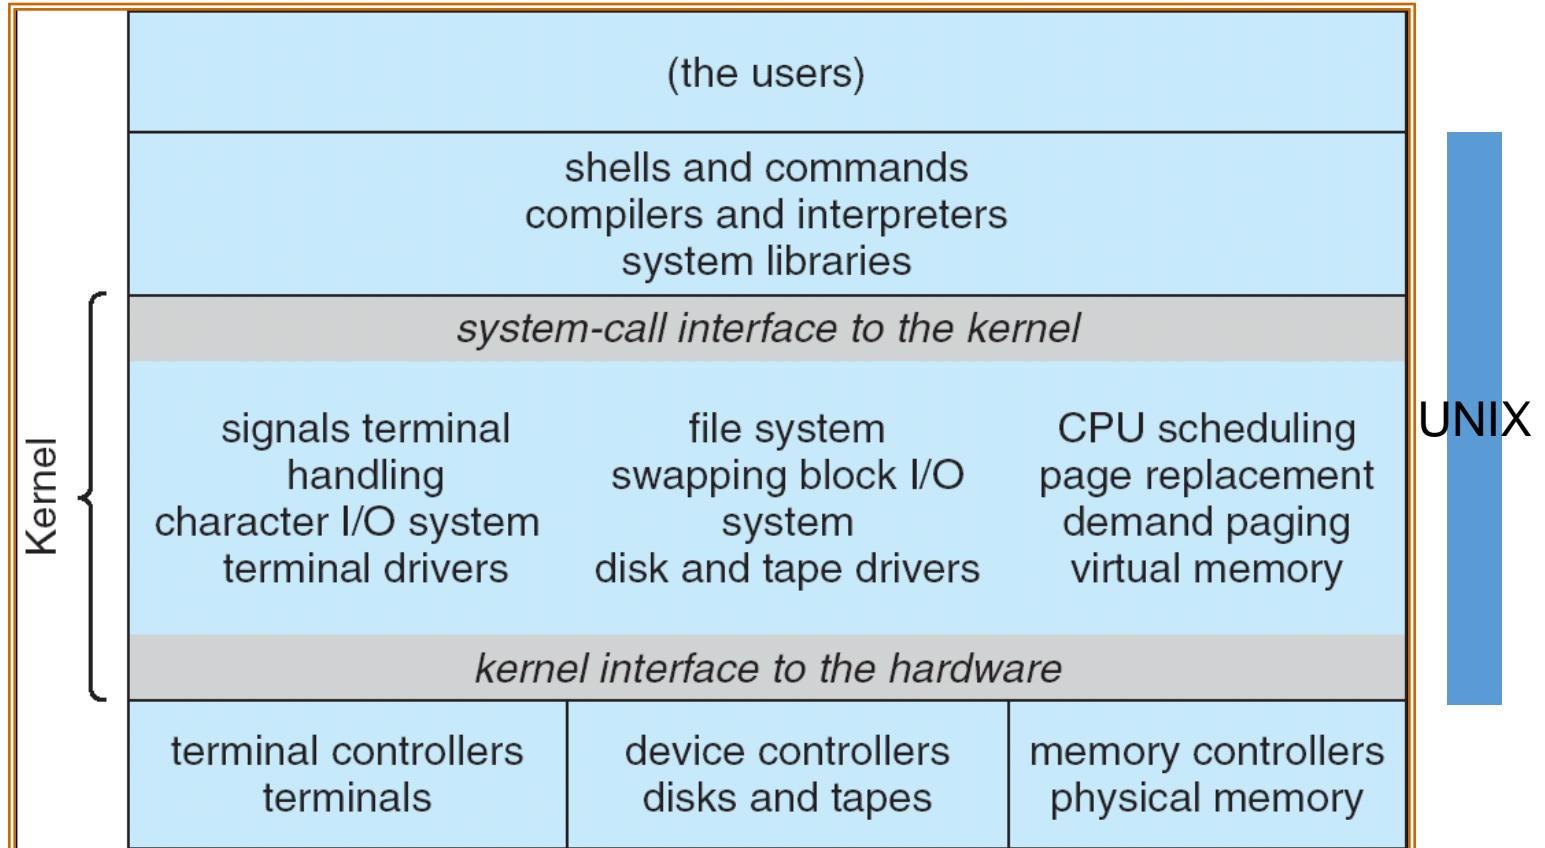
- A real-time, embedded operating systems
- No protection; application + OS in a single image
- Installing new applications is not possible



UNIX--monolithic

- UNIX – limited by hardware functionality, the original UNIX operating system had “limited” structuring.
The UNIX OS consists of two separable parts:
 1. Systems programs
 - binutils + coreutils: ls, mv, cp, etc
 2. The kernel
 - Consists of everything below the system-call interface and above the physical hardware
 - Provides the file system, CPU scheduling, memory management, and other operating-system functions; a large number of functions for one level

UNIX System Structure



UN*X is, of course, a huge monolith operating system
Separation of user space from kernel space

Modules

- Most modern operating systems implement kernel modules
- Common interfaces of a Linux kernel module
 - Initialize
 - Clean up
 - Read (char)
 - Write (char)
 - Read (block)
 - Write (block)

Case Study: Linux Module

```
#include <linux/kernel.h> /* header file for structure pr_info */
#include <linux/init.h>
#include <linux/module.h> /* header file for all modules */
#include <linux/version.h>

MODULE_DESCRIPTION("Hello world !!");
MODULE_AUTHOR("John Doe");
MODULE_LICENSE("GPL");

static int __init hello_init(void)
{
    pr_info("Hello, world\n");
    pr_info("The process is \\"%s\\" (pid %i)\n", current->comm, current->pid);
    return 0;
}

static void __exit hello_exit(void)
{
    printk(KERN_INFO "Goodbye\n");
}

module_init(hello_init);
module_exit(hello_exit);
```

Source:

<http://blog.wu-boy.com/2010/06/linux-kernel-driver-%E6%92%B0%E5%AF%AB%E7%B0%A1%E5%96%AE-hello-world-module-part-1/>

Case Study: Linux Module

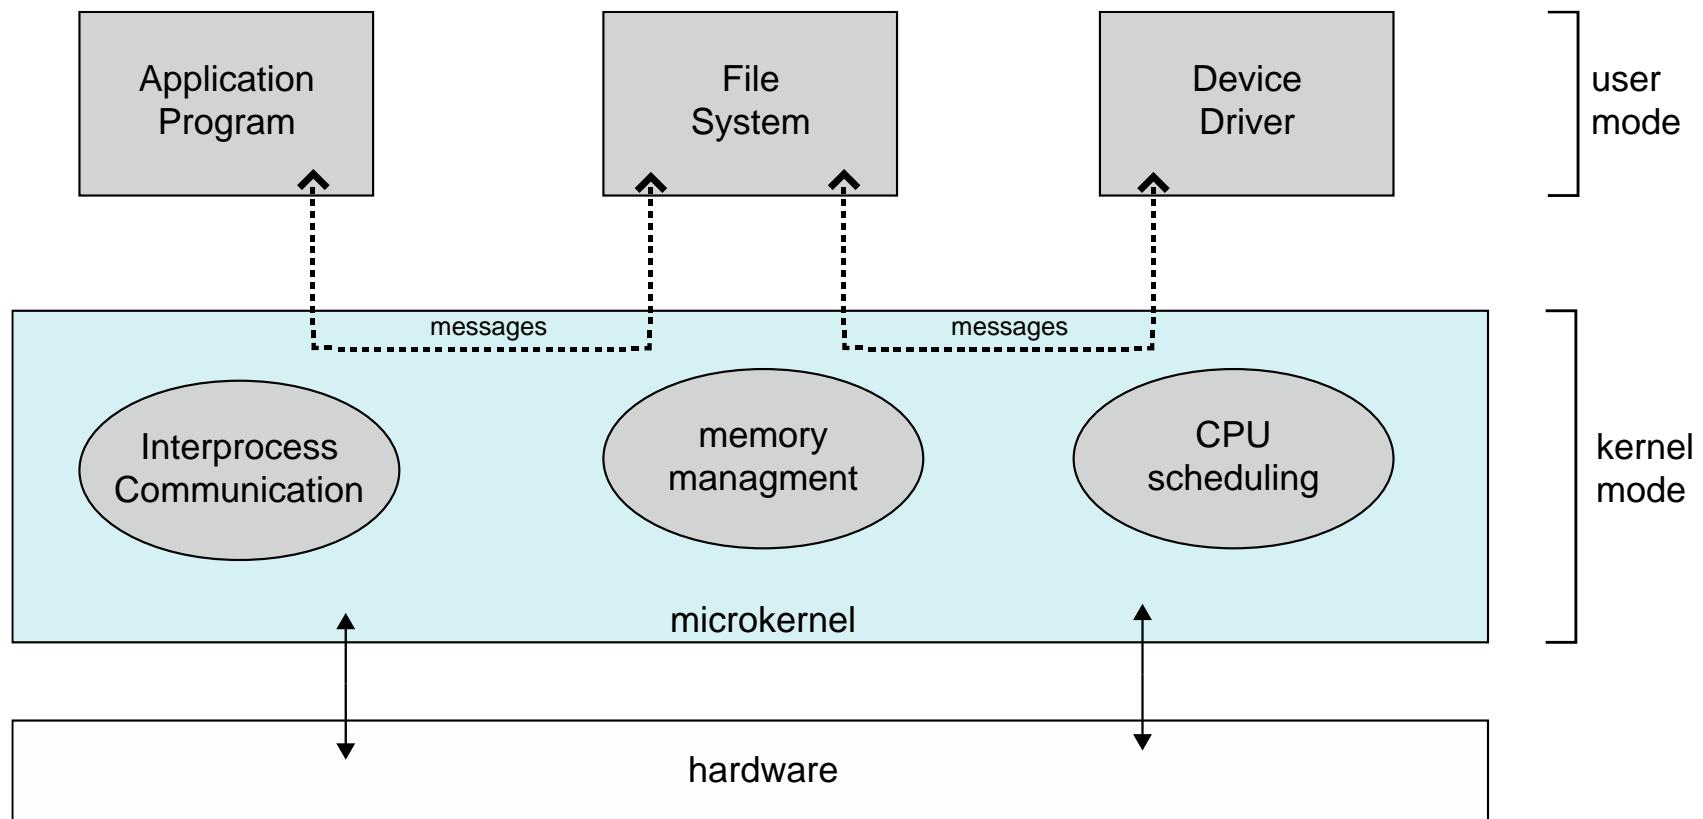
- Insert & init the module
 - insmod ./hello.ko
- Remove & clean up the module
 - rmmod ./hello.ko

```
Jun 21 11:50:00 cvss5 kernel: [8381603.818051] The process is "insmod" (pid 30560)
Jun 21 11:50:08 cvss5 kernel: [8381612.335386] Goodbye
Jun 21 12:07:13 cvss5 rsyslogd: [origin software="rsyslogd" swVersion="4.2.0" x-pid="com"]
Jun 21 12:07:13 cvss5 rsyslogd: [origin software="rsyslogd" swVersion="4.2.0" x-pid="com"]
Jun 21 14:55:23 cvss5 kernel: [8392723.612597] Hello, world
Jun 21 14:55:23 cvss5 kernel: [8392723.612601] The process is "insmod" (pid 10072)
Jun 21 14:55:37 cvss5 kernel: [8392737.604360] Goodbye
Jun 21 15:05:18 cvss5 kernel: [8393318.795982] Hello, world
Jun 21 15:05:18 cvss5 kernel: [8393318.795985] The process is "insmod" (pid 13127)
Jun 21 15:05:25 cvss5 kernel: [8393325.537903] Goodbye
```

Microkernel System Structure

- Moves as much from the kernel into “user” space
- Communication takes place between user modules (processes, specifically) using **message passing**
- Benefits:
 - Easier to **extend** a microkernel (by adding user-mode modules)
 - Easier to **port** the operating system to new architectures
 - More **reliable and secure** (less code is running in kernel mode)
- Detriments:
 - Performance overhead of user space to kernel space communication and user-kernel mode switches
 - What happened to Windows NT 3.5?

Microkernel System Structure

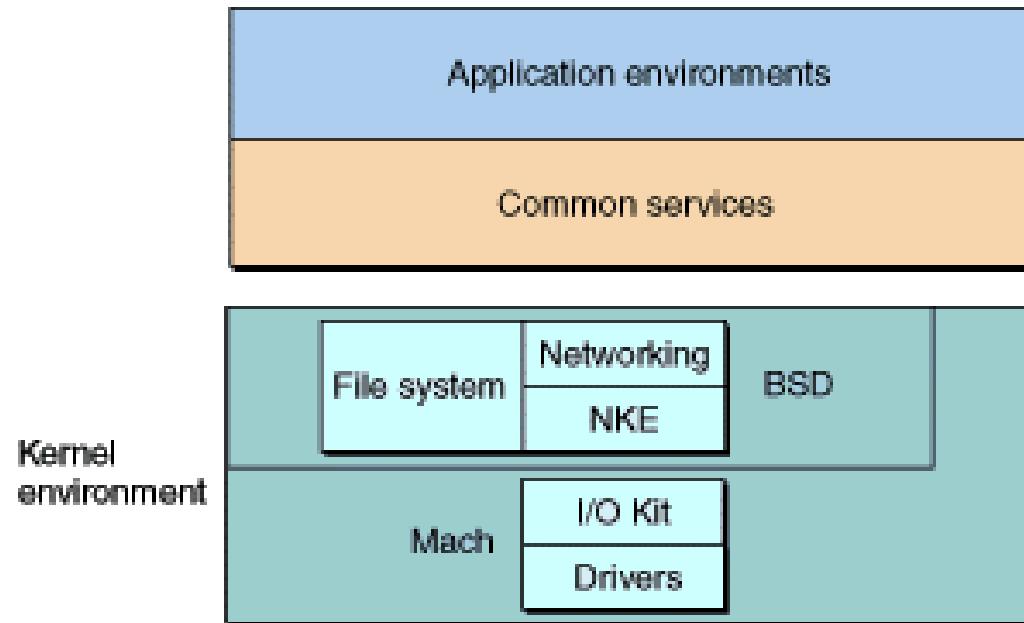


The Famous Tanenbaum–Torvalds Debate

- "*I'm doing a (free) operating system (just a hobby, won't be big and professional like gnu) for 386(486) AT clones.*"
-- Linus Torvalds
- "*LINUX is a monolithic style system. This is a giant step back into the 1970s*"
-- Andrew Tanenbaum
- [Wiki entry](#)

Mac OS X Structure

- Mach (μ -kernel): memory management, RPC, IPC, message passing, thread scheduling
- BSD: networking, file systems, POSIX APIs



Google Fuchsia

- A new operating system developed by Google, based on the Zircon microkernel
- Reportedly designed for IoT devices



Flutter (GUI framework) + Dart (language) → GUI
Fuchsia → system services
Zircon → microkernel

Quiz

What are advantages of the micro-kernel approach?

1. Extensibility
2. Robustness
3. Efficiency
4. Security

Summary: Microkernel

- Provide “a minimal set of kernel primitives”, leave anything else to the user space
- Pros
 - Robust, extensible, secure, portable
- Cons
 - Frequent mode switches/context switches
 - High message-passing overhead

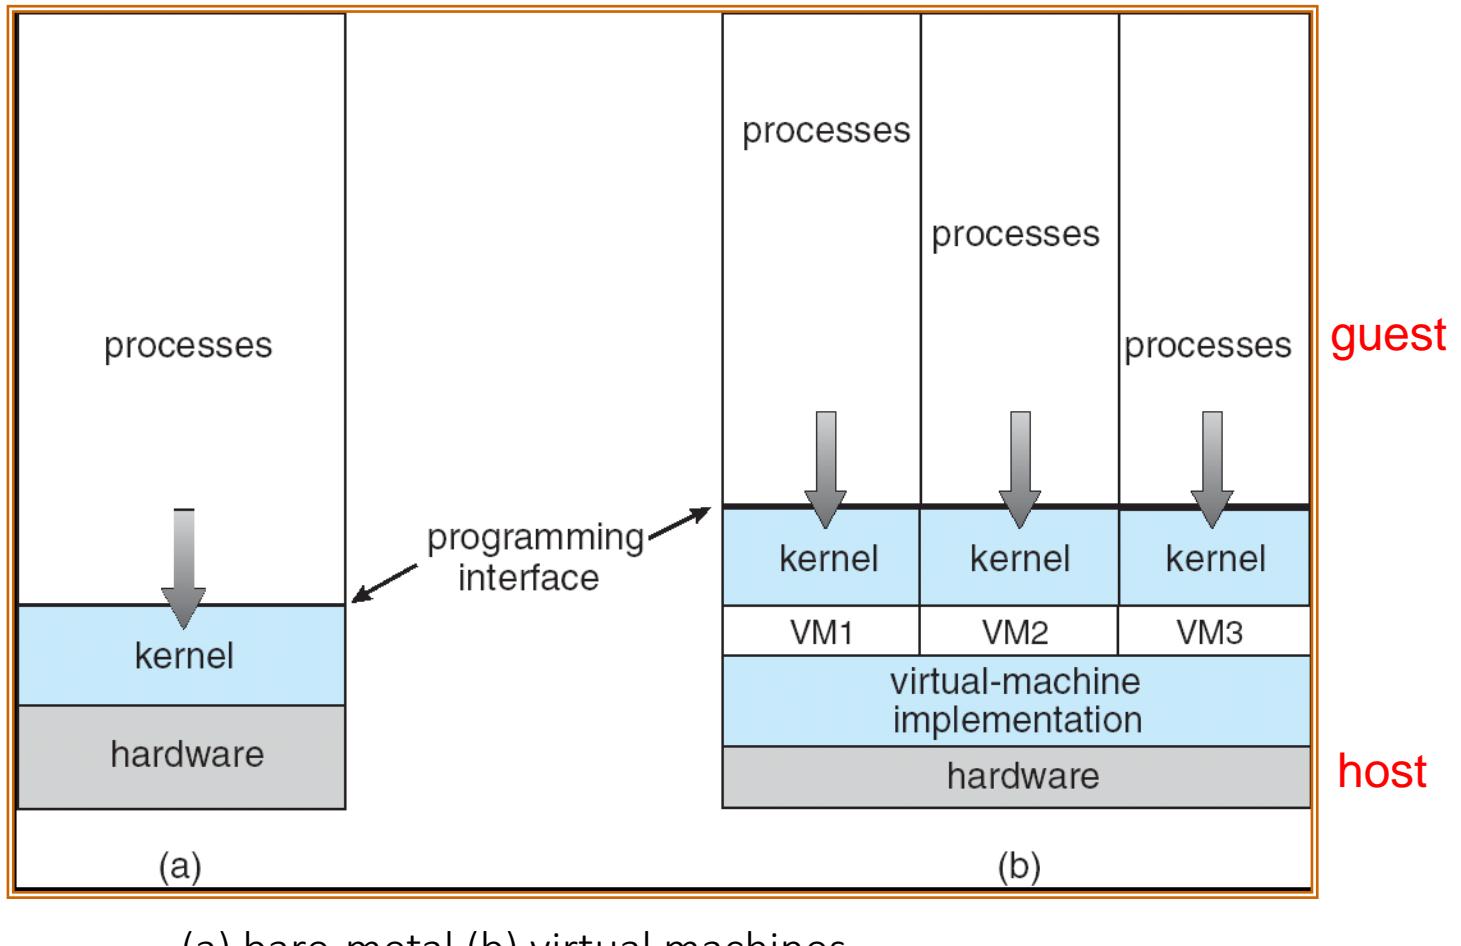
Virtual Machines

- A virtual machine takes the layered approach to its logical conclusion. It treats hardware and the operating system kernel as though they were all hardware
- A virtual machine provides an interface identical to the underlying bare hardware
- The operating system creates the illusion of multiple processes, each executing on its own processor with its own (virtual) memory
- Virtual machine is not a new concept. It has been developed in 197x. VM again become popular because of cloud-based computation
- **Remark:** Different levels of virtualization exist, e.g., system call (Wine), namespace (docker), etc. Here, we focus on the traditional host-guest full system virtualization.

Virtual Machines (Cont.)

- The resources of the physical computer are shared to create the virtual machines
 - CPU scheduling can create the appearance that users have their own processor
 - Spooling and a file system can provide virtual card readers and virtual line printers
 - A normal user time-sharing terminal serves as the virtual machine operator's console

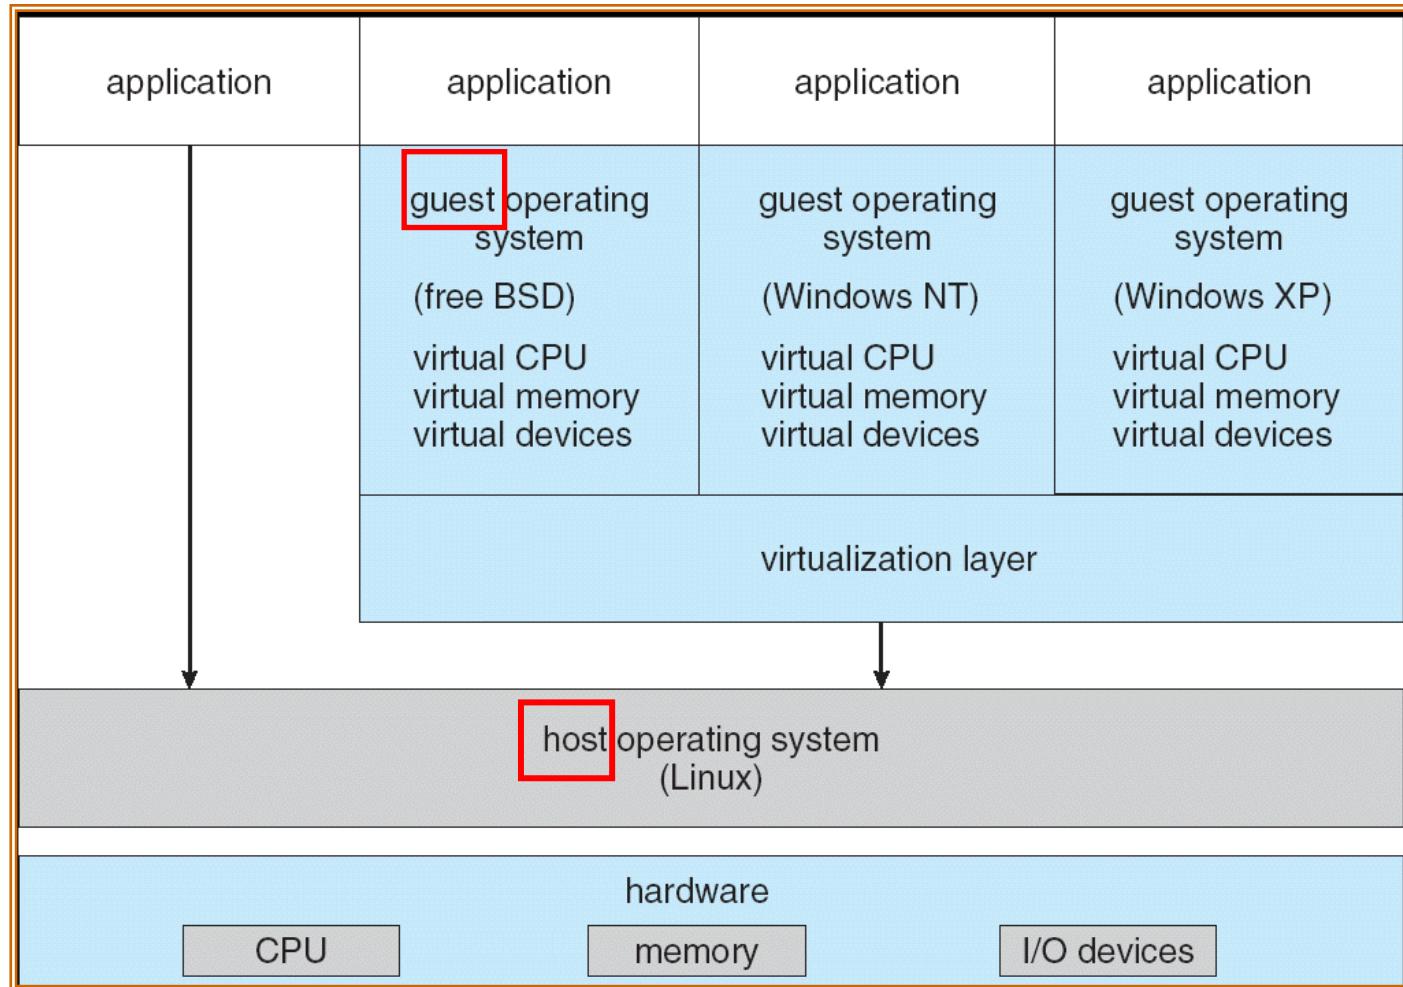
Virtual Machines (Cont.)



Virtual Machines (Cont.)

- The virtual-machine concept provides **complete protection of system resources** since each virtual machine is isolated from all other virtual machines. This isolation, however, permits no direct sharing of resources.
- A virtual-machine system is a **perfect vehicle for operating-systems research and development**. System development is done on the virtual machine, instead of on a physical machine and so does not disrupt normal system operation.
- The virtual machine concept is **difficult** to implement due to the effort required to provide an exact duplicate to the underlying machine

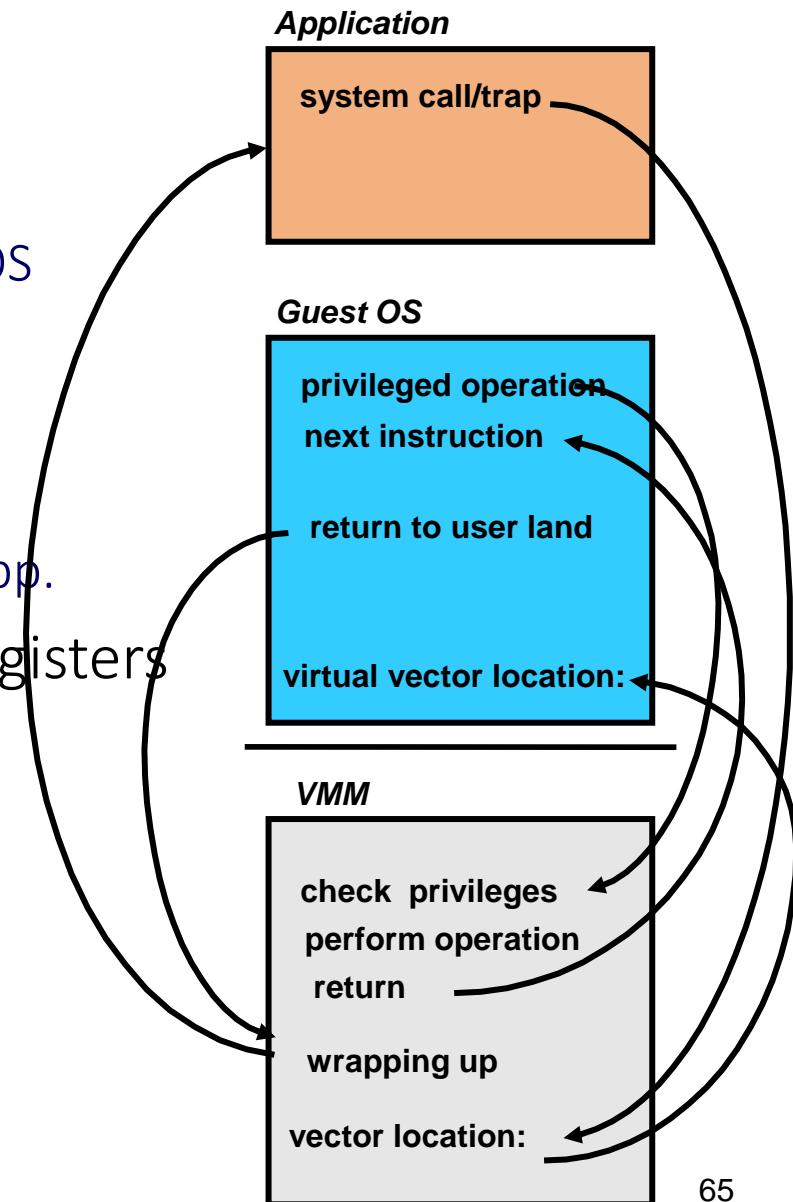
VMware Architecture



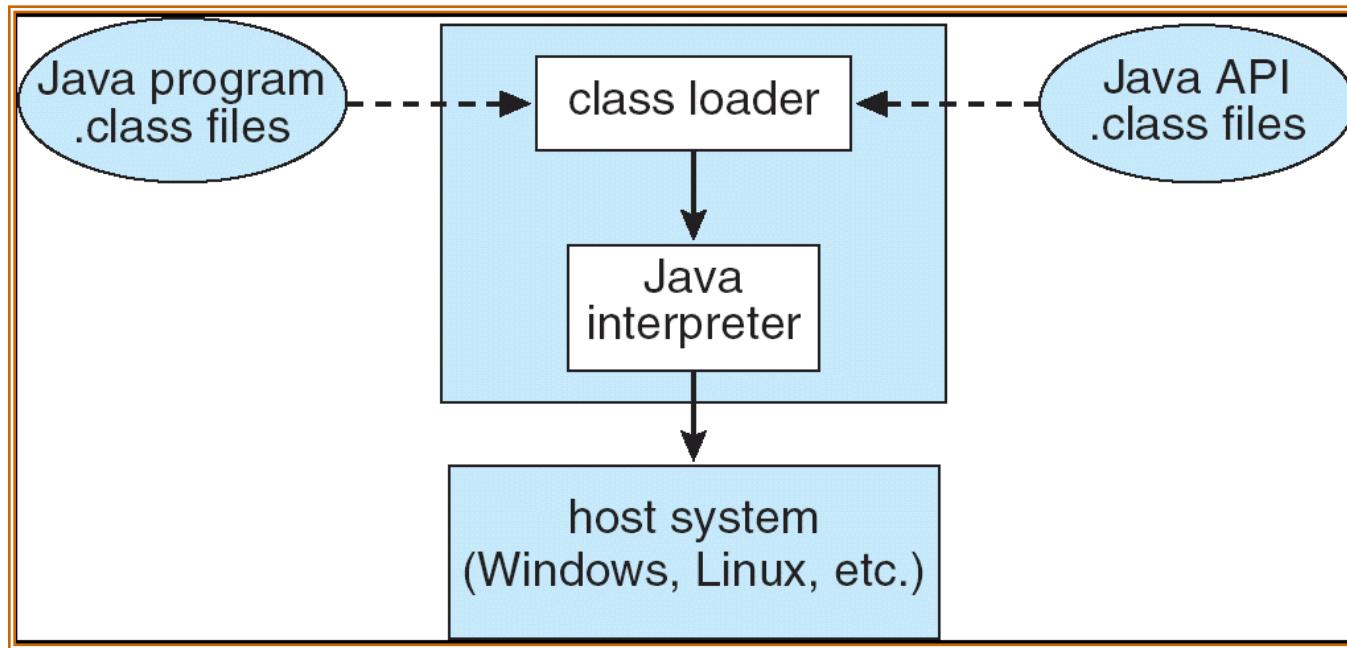
Native execution

Processor Management/Protection

- Traps and interrupts (& sys calls)
 - Transfer to VMM
 - VMM determines appropriate Guest OS
 - VMM transfers to Guest OS
- Guest OS “return” to user app.
 - Transfer to VMM
 - VMM bounces return back to Guest app.
- Read/Write of protected control registers
 - Trap to VMM
 - VMM reads/modifies guest copy
 - May modify shadow copy
 - Returns to Guest



The Java Virtual Machine



Java programs: the source code

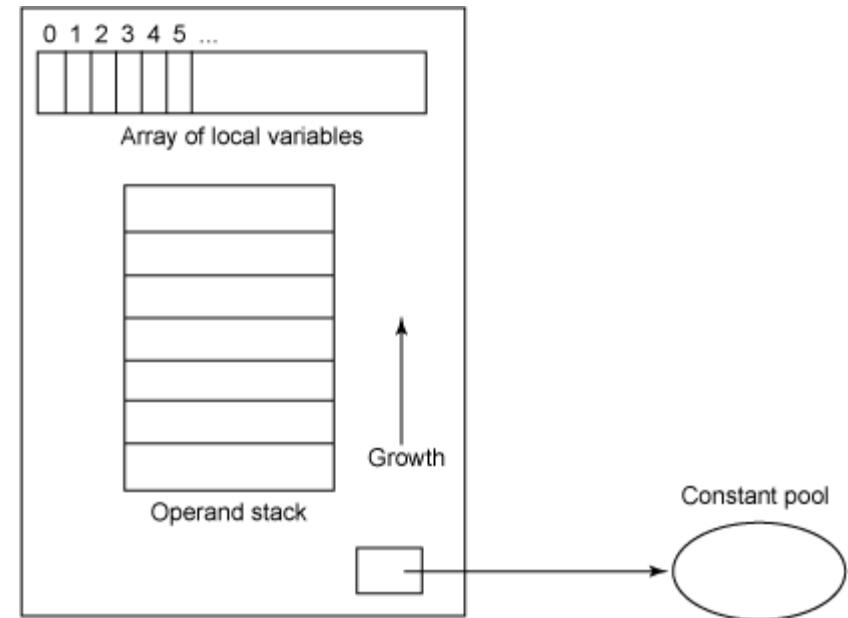
Byte code: the compiled binary for JVM

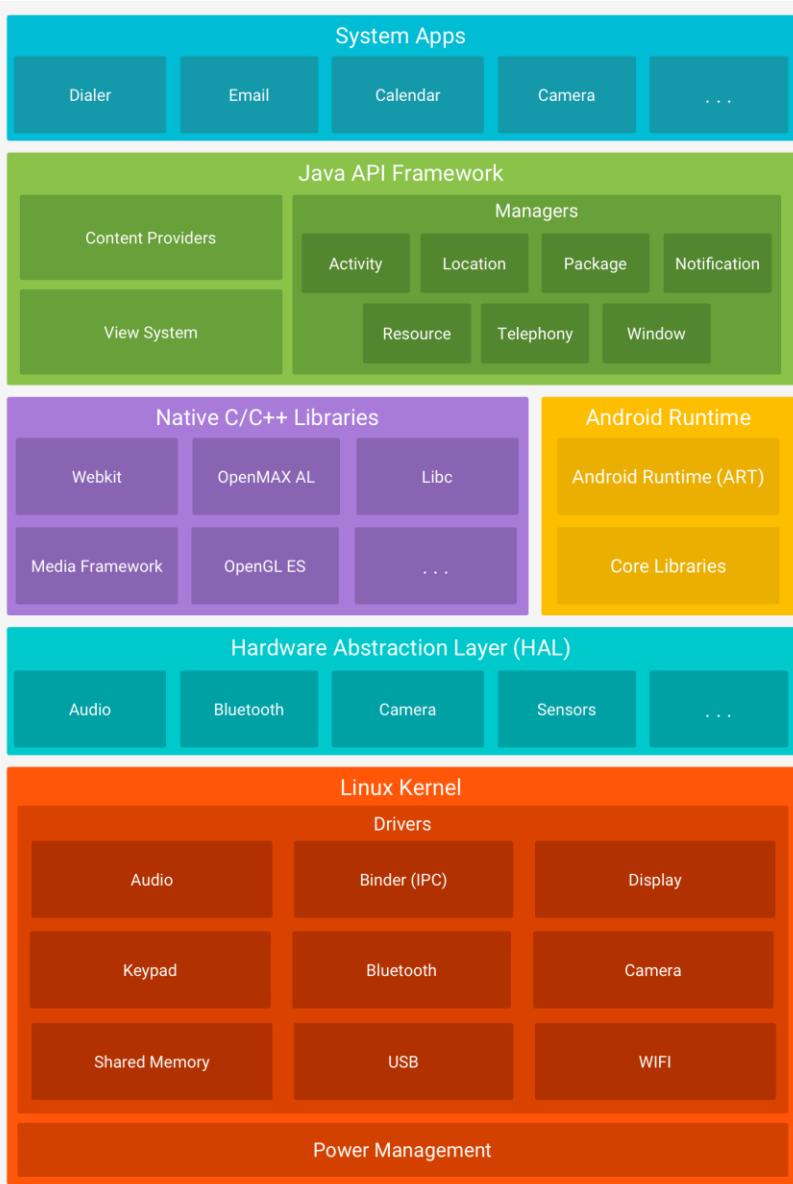
JVM or java runtime: a hardware-independent virtual machine

*Non-native execution

Java Bytecode

iload_1
iload_2
iadd
istore_3





Android Structure

- **Java API Framework:** The entire feature-set of the Android OS is available to you through APIs written in the Java language.
- ART is written to run multiple virtual machines on low-memory devices by executing DEX files
- The hardware abstraction layer (HAL) provides standard interfaces that expose device hardware capabilities to the higher-level Java API framework.

Quiz

The virtual machine approach is suitable to which one(s) of the following scenarios?

1. OS development
2. Cloud computing
3. Performance-critical gaming
4. Writing an application for heterogeneous hardware platforms

Summary: Virtual Machines

- Virtualizes hardware
- Pros
 - Guest operating systems run without modifications
 - Perfect resource partition and fault isolation
 - Resource sharing among VMs
- Cons
 - Inefficient mapping between emulated hardware and the underlying hardware
 - Hard to implement hardware virtualization

Review: Different OS Structures

- Simple structure
 - Pros: simple, cons: poorly structured
- Monolithic
 - Pros: efficient, cons: hard to scale
- Microkernel
 - Pros: robust and scalable, cons: inefficient
- Virtual machine
 - Pros: perfect resource isolation, cons: inefficient hardware emulation

Fun Fact

What is the most popular operating system for personal computers?

- It might be MINIX. It is part of INTEL ME, integrated in the motherboard chipset

End of Chapter 2

Review Questions

1. Establish the mapping between fopen(), open(), & int 80h #5 and libc call, system call, & POSIX API
2. Why system calls are implemented as software interrupts (traps) but not standard function calls?
3. Explain which part of CPU scheduling is a mechanism, and which is a policy
4. Why monolithic kernels are more successful in real life computer systems?
5. In embedded systems, there is no separation between user mode and kernel mode. Why?
6. How a system call from the guest OS is handled?
7. Why micro kernels are considered more secure?
8. Briefly survey Wine (system call emulation), Docker (name-space), Xen (para-virtualization), and QEMU (full virtualization)